

03 DATABASE MANAGEMENT SYSTEM

Module - I

Database system architecture; Data Abstraction, Data Independence, Three-Schema Architecture
 Data models: Entity-relationship model, network model, relational and object oriented data models, integrity constraints.

Module - II

Relation Query Languages: Relational Algebra, SQL, Data Definition Language (DDL),
 Data Manipulation Language (DML), Tuple and Domain Relational Calculus.

Relational Database design: domain and data dependency, Armstrong's axioms,
 Normal forms, Dependency preservation, Lossless design.

Module - III

Query processing and optimization: Evaluation of Relational Algebra Expression,
 Query Equivalence, Join strategies, Query Optimization Algorithms.

Module - IV

Transaction processing: Concurrency control, ACID property, Serializability of scheduling, Locking and timestamp based schedulers, Multi-version and optimistic concurrency control schemes, Database recovery.

Module - V

Advanced Topics: (Introduction to concepts only) Object oriented ~~data~~ and object relational databases, logical databases, Distributed databases, Data warehousing and data mining.

L-1 (27.07.2023) (11:00 - 12:00) DBMS

Objectives:

- Basic concepts, modeling, architecture
- Design aspects of database
- Use of DDL & DML for query purpose
- Query evaluation and optimization
- Basic issues & concept - Transaction management.

- Remembering
 - Understanding
 - Applying
 - Evaluating
 - ↳ Creating
- (Bloom's Taxonomy)

* State of inconsistency, failed transaction, recovery has to be done.

The log file records each activity. Helps in recovery management.

* Redundant Data: data being referred in many places/multiple places.

* Vendors of Database: MySQL, Oracle, MS SQL, MongoDB etc.
 DBMS is an application software that facilitates storage, access & different activities with the DB.

☒ Textbook: 6e or 7e Silberschatz, Korth, Sudarshan (Database System Concepts, 7e)

☒ Enroll IIT KGP SWAYAM online course - Partha Paitam Das. (www.db-book.com)

① Railway reservation system: Robust Database system

* shopping, social media platforms: use multimedia data (no specific structure)

→ NOSQL (RDBMS is not sufficient) → Referred to as Big Data

File system DB

↳ Data capturing & handling, similar to XML/HTML (use of tags)

↳ Modern Database Applications.

↳ Data Redundancy
 ↳ Inconsistency
 ↳ Difficulties in Access

L-2 (28.07.2023) (11:00 - 12:00 p.m.)

Database Management System is an application software that implements database and its applications.

① 2-tier and 3-tier architecture

↓
client, application layer, db layer

Client/user sit on the frontend side

Data storage & operations - backend application server

↳ Data Models:

We use data models to describe certain levels of data abstractions.

→ Data models are the conceptual tools, used to describe certain levels of data abstraction.
↳ high level description of data.

Represent: Data, Data relationships, Data semantics, Data constraints

→ Representational Data Models: no detailed physical description required

→ Physical Data Model: physical description of data.

• Relational model

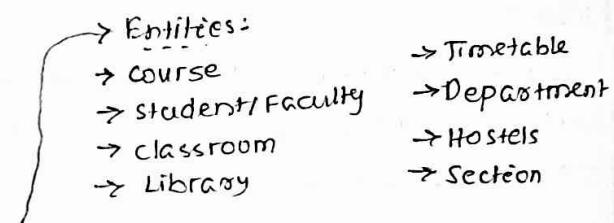
- Entity-Relationship model (ER-model)
- object-based data model (object-oriented & object-relational)
- Semi-structured data model (XML)

Older Data Models:

- Network model (form of a graph)
- Hierarchical model (form of a tree)

ER-model

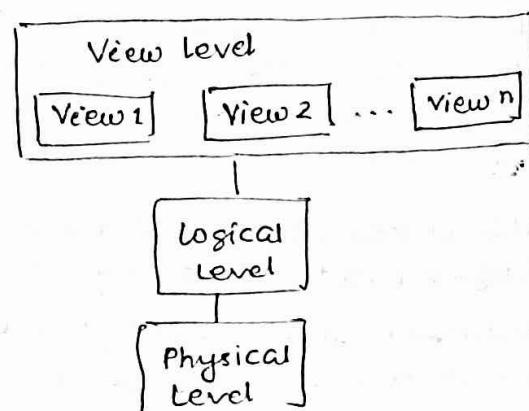
- provides concepts of entity, relationships & attributes.
↳ identifying characteristics.



• E.g.: University Database

• Relationship: Interaction among student & enrolled course

Isolated access for different kind of user (conceptual level)



↳ Relational Data Model:

Ted Codd: Turing Award 1981

• All the data is stored in various tables.

columns

	ID	name	dept-name	salary
t1	222	X	Physics	50,000
t2	213	Y	Finance	95,000

rows

• The instructor Table

- Each row is called as Tuple/Instance
- columns represent attributes.

↳ Instances and Schemas:

Instructor (ID, name, dept-name, salary)

↳ is a schema

↳ dept (dept-name, building, budget)

Schema is analogous to variable

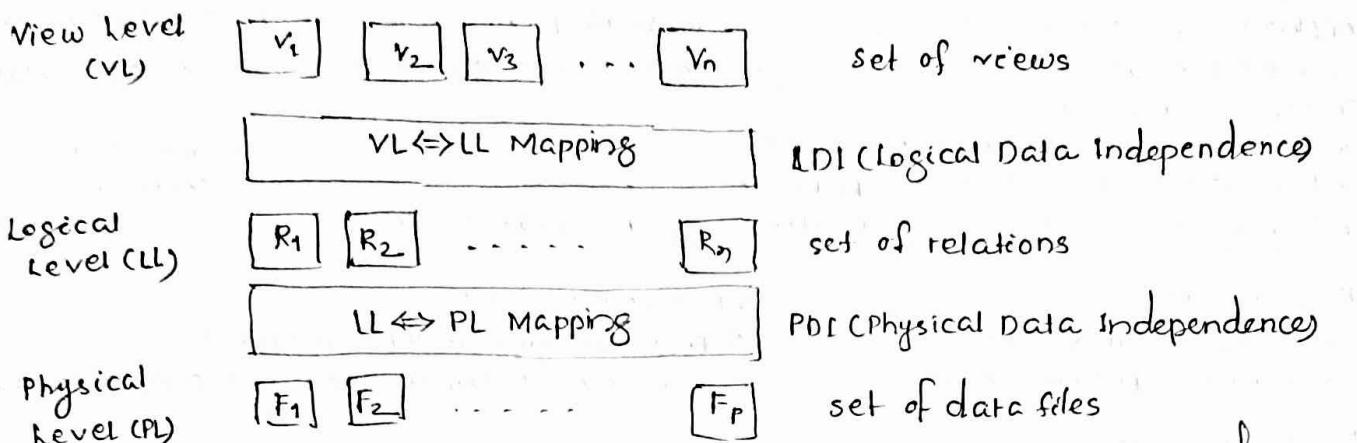
Instance is analogous to value assigned to the variable

→ In database, data and schema are stored separately.

→ Schema refers to table name, attribute name, data types and constraints.

- character as well as number - VARCHAR
- numbers, strings ..

- Primary Key, Foreign key are specified as constraints while defining the schema.
- Database Definition:
 - ↳ Setting up the skeletal structure
- Database loading/populating:
 - ↳ Storing data



- The whole database is not projected to all. A portion is projected to individual view.

• L-3 (03.05.2023) (11:00-12:00 p.m.) (12:00-12:30 p.m.)

- The development process of a database system: 3 level process

Step:01 Gathering/ Collecting the requirements.

Step:02 Conversion of data models into representational model (ER/Object-oriented/Semi-structured)
identification of model

Step:03 Write the application program using general purpose programming language.

① In database perspective, requirements are of two types:

- Data model requirements
- Functional requirements

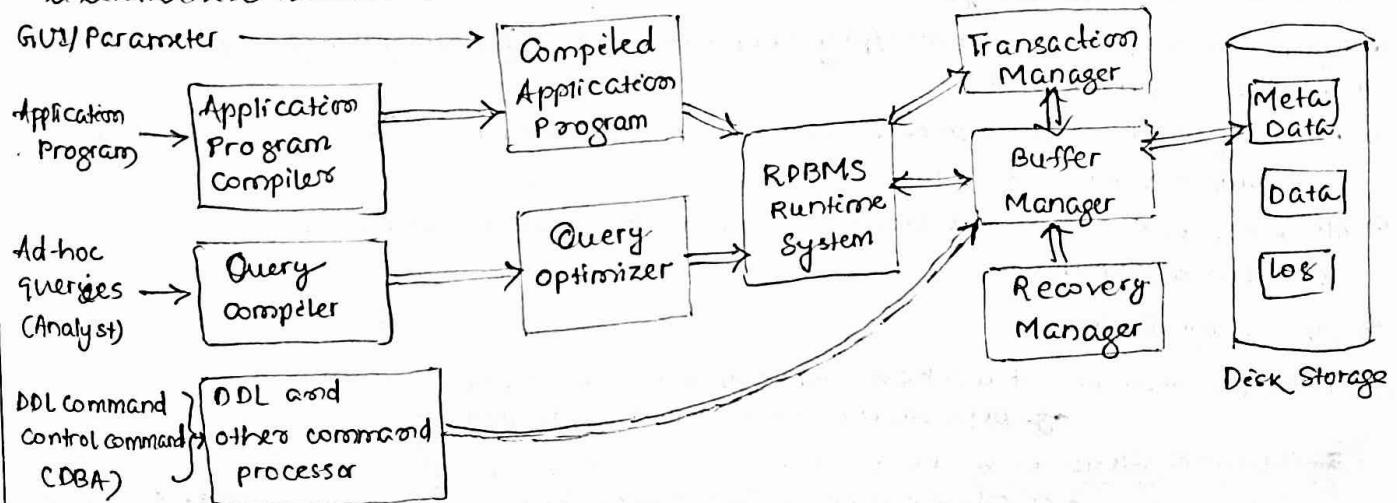
embedded into
SQL Queries

② Through the application program, we access the database values, through some variables.

③ Embedded SQL, using SQL commands within the high level language to interact with the database and carry out necessary functionalities.

[B] Navathe Book

- Architecture of RDBMS System:



- ① Data dictionary contains definition of table, mappings & schema.
- ② Meta data stores / contains the data dictionary.
- ③ Log contains activity record, old value as well as new values for any change in database value.
↳ used for recovery purpose.
- ④ DDL : Data Definition Language, DBA : Database Administrator
Commands for handling authorization is also controlled by DBA.
- ⑤ Query optimizer selects the appropriate execution plan. Also optimizes join operations.
- ⑥ Application program compiler separates the embedded SQL commands. Application program is compiled separately. SQL commands are compiled. And then both are integrated. (3 tasks)
- ⑦ Transaction Manager
 - Tracks the start & end of a transaction
 - Records the changes in log protocols
 - Emphasizes concurrency control, (Process synchronization. - OS)
↳ enforces
 - Time stamp Protocol
- ⑧ Buffer Manager
 - ↳ manages the disk space.
 - ↳ implements paging mechanism.
- ⑨ Recovery manager
 - ↳ Performs some undo/redo operation.
 - ↳ Takes control, brings system to a consistent state.

★ 2 tier vs 3 tier..

Levels of abstraction..

↳ Different User Roles:

~~~~~

#### ① DBA (Database Administrator)

- configure database
- write DDLs
- physical layout
- creating data dictionary
- design logical schema
- grant & revoke permissions  
(use command line interfaces)

↳ ER Model:

~~~~~

→ Proposed by Peter P. Chen (1970)

→ Trying to establish relationship between the entities.

→ Graphical representation: ER-Diagram

ER-Diagram convert into Schema

→ Try to identify the attributes for individual entity.
(what describes that entity).

② Entity is a thing of independent/physical existence, which is distinguishable.

Individual database contexts.

③ An entity set is a collection of entities carrying all same properties (characteristics).

e.g.: student entity set is a collection of all student entities.

④ Attribute: Entity set is described by attribute, that carries some value.

e.g.: Roll no. in STUDENT entity set.

⑤ Types of attributes:

i. Simple attribute : it will have atomic or indivisible values.

e.g.: Department name cannot be divided.

ii. Composite attribute: We have several components in the value

e.g.: Qualification → (start year, end year, institution, name of degree)

↳ the subparts of qualification

② Data Entry Operators

- Provided with GUI through which they interact with DB.
- Feed the data and invoke particular operations.

③ Sophisticated User

- Analysts
- specify queries to get the data (SQL Query)

e.g.: Department Data Base

④ Entities: Department, Faculty, Student, Course, Enrollment..

⑤ Relation: Faculty offers a course

teaching

Student belongs to department.

Student enrolled in a course.

3. Derived attribute: value depends on some other attribute value.
e.g.: age - can be derived from DOB.

4. Single-valued attribute: Some attributes can only take a single value.
e.g.: place of birth.

5. Multi-valued attribute: (email ID may have more than one values)

• simple single-valued attribute, simple multi-valued attribute

• composite single-valued attribute, composite multi-valued attribute

Q How exactly the types of attributes are represented? : Diagrammatic Notation

→ Entity - Rectangle

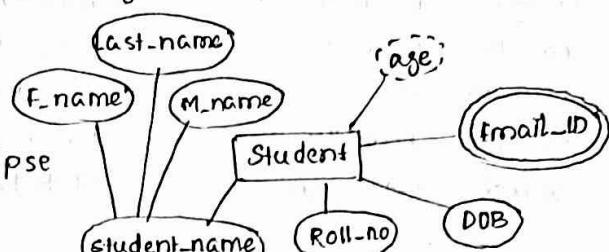
→ Attribute - Ellipse connected to rectangle

→ Multivalued attributes - double ellipse

→ Composite attributes - ellipse connected to rectangle

→ Derived attributes - dashed ellipse

Example: Entity: Student



Q Representation of an entity in ER Diagram.

→ L-4 (04.08.2023) (11:00-12:00 P.M.)

↳ Domain of attributes:

• Each attribute takes the value from a set called 'domain'.

• e.g.: age domain of students.

• Domain of composite attribute is the cross-product of domains of sub-attributes.

• Domain of multivalued attributes is the set of sub-set values from the basic domain.

↳ What is a key:

• Key is an attribute or combination of attributes which is used to uniquely identify a particular tuple of an entity set.

• For a particular entity set, we have more than one key(s) possible.

• Candidate key set: set of keys.

• The minimal candidate key is referred to as the 'primary key.'

{ Designer evaluates meaning of individual attributes

forms the candidate key set, finds the minimal of the keys: primary key.

↳ Relationship:

• The association among the entities is referred to as relationship.



(diamond: represents relation)

Relationship set

○ enrollment ⊆ student × Course

○ (s, c) ⊆ enrollment i.e., student s has enrolled course c.

↳ Degree of Relationship:

(1) Binary Relationship

(2) Ternary Relationship

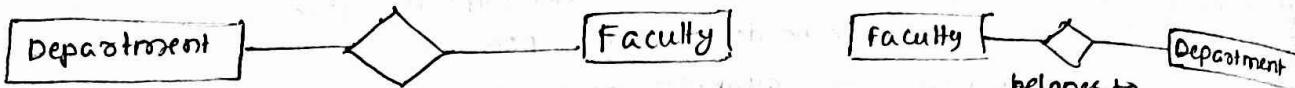
(3) n-ary Relationship

• We say the no. of participating entities as degree.

• In ER-model, binary relationship is mostly used.

↳ Examples of binary Relationships:

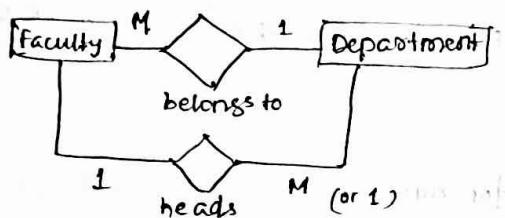




- Department has faculty
- Faculty belongs to the department

Entities

			<u>Relationship</u>
student	course	—	enrollment
student	Department	—	belongs to
faculty	Department	—	belongs to
faculty	Course	—	offers / teaching
student	Faculty	—	is mentored by / mentorship
course	Classroom	—	allotted / assigned



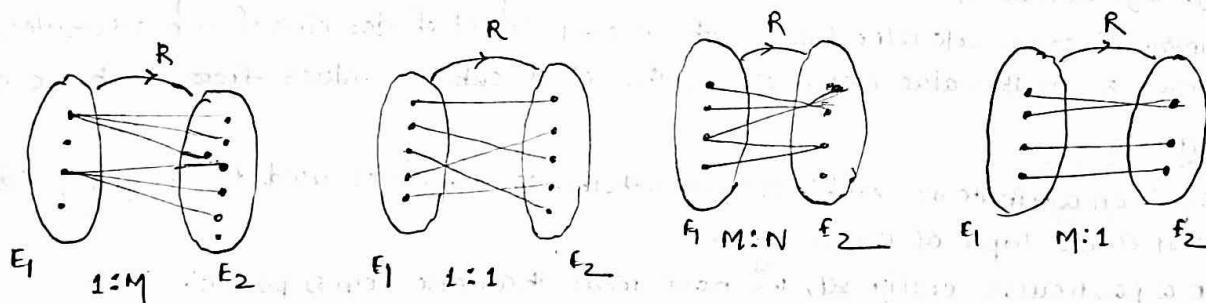
② Cardinality Ratio

(Set of Relations)



→ cardinality ratio in binary relationship has four possibilities:

- 1:1 (one-one)
- 1:M (one-many)
- M:1 (many-one)
- M:N (many-many)



Examples:

One-to-one:



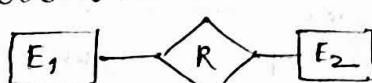
Many-to-one/ one-to-many:



Many-to-many:



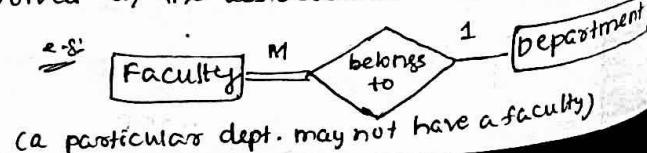
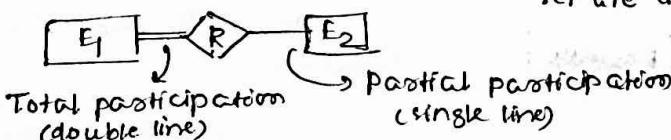
Participation Constraints in a Relation:



① In case of total participation all the entity sets should participate in the relationship.

- T.P. (Total Participation)
- P.P. (Partial Participation)

② In partial participation, not all entities in the set are involved in the association.



1-5 (05.08.2023) (10:00 - 11:00 a.m)

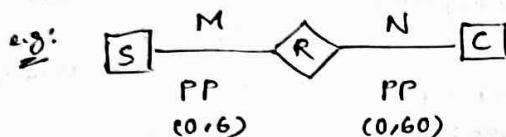
M-Max: (M, n)

- Need to be specified along the line from entity to relationship.
- 'M' is the minimum number of times a particular entity must appear in the relationship at any point of time.

In case of partial participation, $m=0$

But in total participation, $m \geq 1$

- 'n' is the maximum number of times a particular entity can appear in the relationship tuple at any point of time.

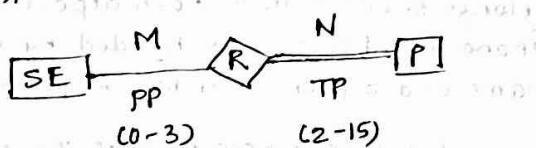


All the students needn't register the course, but they can go maximum upto 6 courses

All the courses needn't be registered by a student, if at all they register, maximum 60 students can take up the course.

e.g: Each and every project requires atleast two members and maximum of fifteen members. All the software engineers needn't be the part of project, if they participate, they can participate in maximum three projects. Find the type of relation and cardinality ratio.

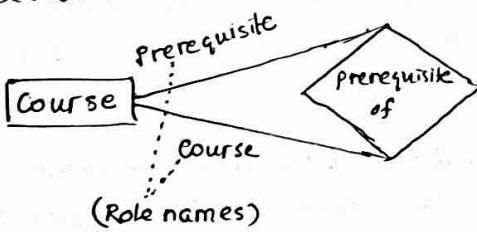
→ It is a binary relationship.
with cardinality ratio M:N
(many-to-many)



Attributes for Relationship Type:

A relationship may have some attributes.

Recursive Relationship:

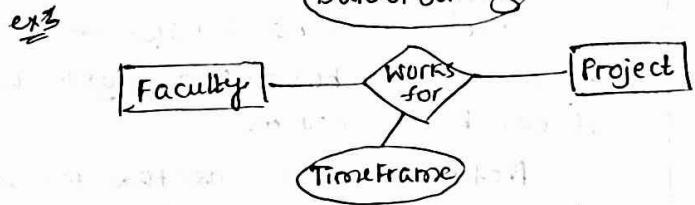
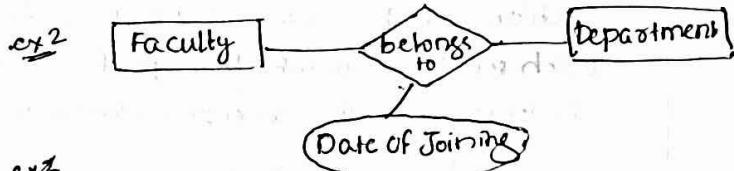


↓ used to specify the exact roles in

which the entity participates in the relationships

↓ are essential in recursive relationships

↓ are not essential in non-recursive relationships



Entity Sets:

→ strong entity sets
→ weak entity sets

• Entities having uniquely identifying attribute ↓ strong
↓ own attribute will be used as key value ↓ entity set

• Attributes of entity are not sufficient ↓ weak
for identifying entity sets uniquely ↓ entity set

↓ The members are dependent on existence of strong entity set
i.e. weak entity set doesn't have any independent existence.

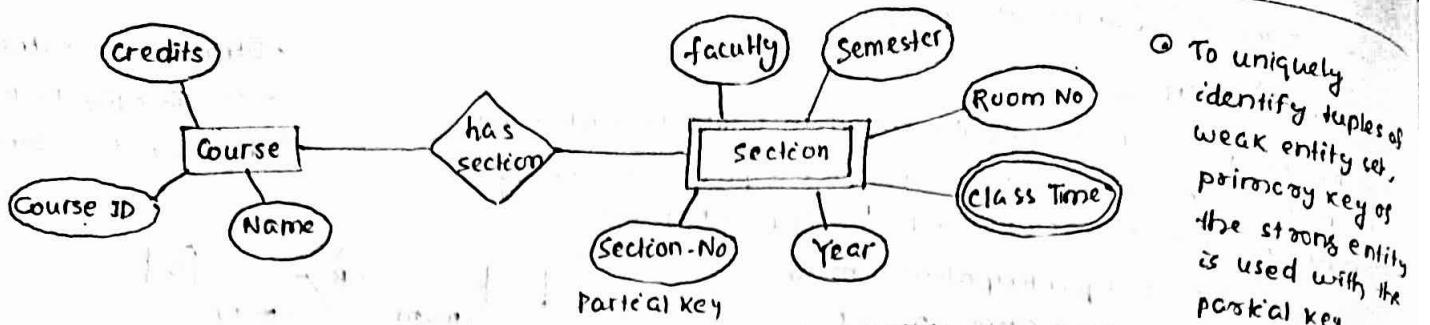
- ↓ Structural constraints:
- cardinality Ratio
- Participation constraint

e.g: Student - strong

Section - weak

DBMS - Prerequisites

- Linear Algebra
- set Theory
- OSA
- Basic Programming



① To uniquely identify tuples of weak entity set, primary key of the strong entity is used with the partial key.

e.g: Equipment (strong), Usage(weak), Relationship: Utility

② Added intentionally to a relation to uniquely identify tuples in the relation — Discriminator (in weak entity sets)
(e.g: token number for transactions)

Real-time simulation
as well as emulation
OPAL-RT system
(70-80 layers)

③ Discriminator: combination of primary key & partial key
(strong) (weak)

clear ER-diagram

& Complete Example for E/R schema:

Data model requirements are given.

Specifications: In an educational institute, there are several departments and each student belongs to one of them. Each department has a unique department number, a name, a location, phone number and is headed by a professor. Professors have a unique employee id, name and a phone number. A professor works for exactly one department.

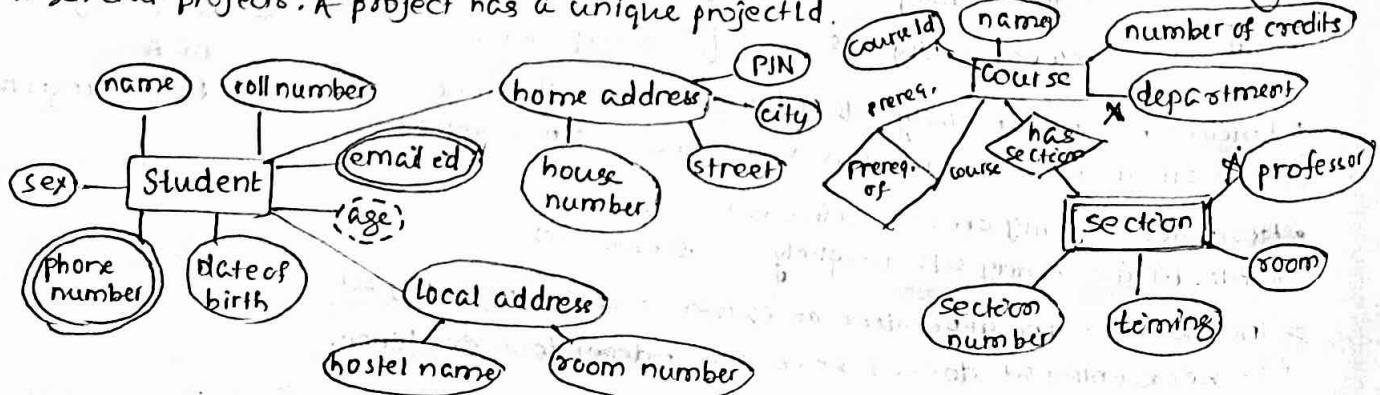
We like to keep track of the following details regarding students: name, unique roll number, sex, phone number, date of birth, age and one or more email addresses. Students have a local address consisting of the hostel name and the room number. They also have home address consisting of house number, street, city and PIN. It is assumed that all students reside in the hostels.

A course taught in a semester of the year is called a section. There can be several sections of the same course in a semester; these are identified by the section number. Each section is taught by a professor and has its own timings and a room to meet. Students enrol for several sections in a semester.

Each course has a name, number of credits and the department that offers it.

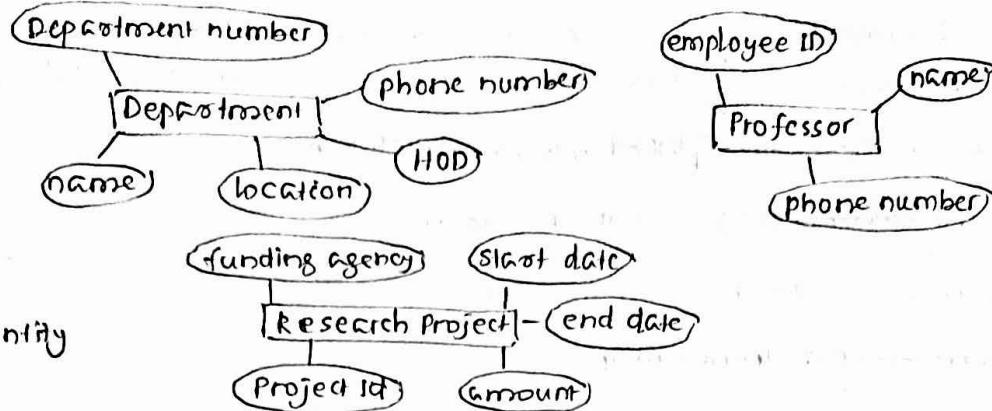
A course may have other courses as pre-requisites i.e., courses to be completed before it can be enrolled in.

Professors also undertake research projects. These are sponsored by funding agencies and have a specific start date, end date and amount of money given. More than one professor can be involved in a project. Also a professor may be simultaneously working on several projects. A project has a unique project id.

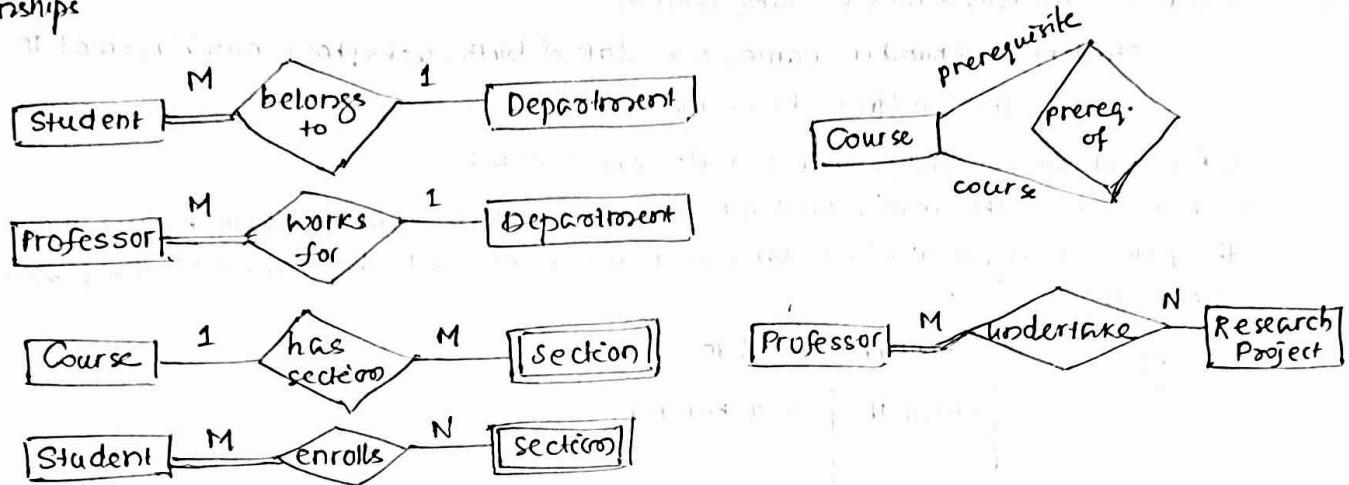


Entity Set

- Department
- Student
- Professor
- Course
- Section -- weak entity
- Research Projects



Relationships



L-6 (10.08.2023) (11:00 - 12:00 PM)

Student

- name
- roll number
- sex
- age
- email id
- date of birth
- local address
- home address

Professor

- professor ID
- name
- phone number

Section

- section ID
- Timing
- classRoom

Department

- Department ID
- name
- location
- phone number
- HOD

Course

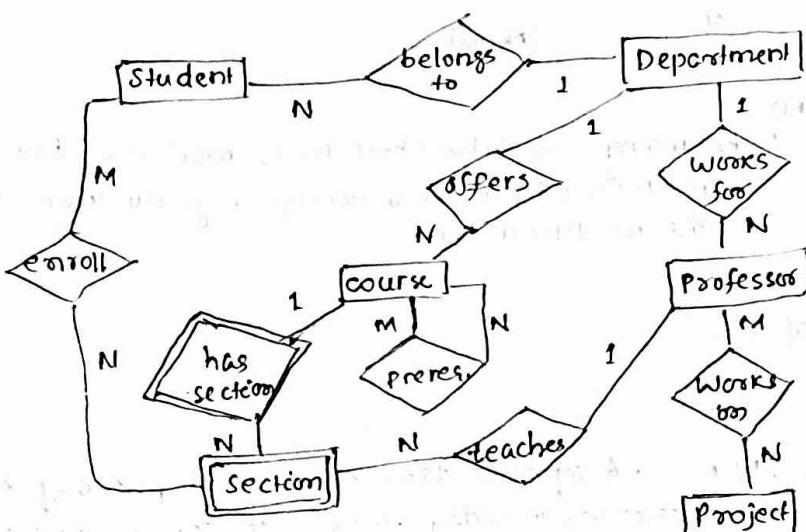
- Course ID
- name
- number of credits

Project

- Project ID
- start date
- end date
- sponsor
- amount



weak relationship



④ While defining the key of a weak entity, we borrow the primary key of the associated strong entity. The combination of this along with the ID of weak entity form discriminator.

→ Ternary relations are splitted into binary relations.

+ Designing different tables for the ER Diagram:

(Conversion of ER-Diagram into tables)

- Strong entity set
- Weak entity set
- Relationship

1. Conversion of a strong entity sets:

→ Consider all attributes as table column.

student (roll-number, name, sex, date-of-birth, age, phone-number, email-ID, local-address, home-address)

→ For each strong entity, separate tables are created.

→ For each of the multivalued attributes, create a separate table having the attributes as the primary key of the base table as a reference and the values of the particular attribute.

e.g:

TABLE email-ID

email-ID	roll-number

→ We can break the composite attribute into simple attributes. Capture sub-attributes separately.

2. Conversion of weak entity:

e.g: Section (course-ID, section-ID, Timing, class-room)

→ Keep all attributes of the entity set along with the primary key of associated strong entity.

3. Relation:

→ One-to-one relationship (1:1)

- Include primary key of one entity set as column; foreign key to the table we want to modify (either A or B)

• If relationship has an attribute, this too is included in the modified table.

A
-- PK(B) x

or

B

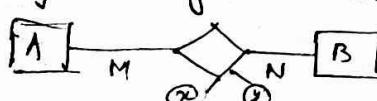
→ One-to-many relationship (1:M)



• The many side table (here B) is modified. Primary key of A is referenced by B as a foreign key. Attributes included in the modified table.

B
-- PK(A) x

→ Many-to-many relationship (M:N)



• A separate table is created, primary key of both tables are included along with the relationship attributes.

A					

B					

PK(A)	PK(B)	x	y

* Creation of Table for the ER-Diagram: (P-09)

• Insertion
Deletion
anomaly.

Student

student_ID/Roll-number	name			sex	date_of_birth	age	local add.	home add.	Dept.
	first	mid	last				h-name	r.no.	h-n st. city PAN ID

Emails

email-ID	Roll_number

Department

Dept-ID	Name	Location	Phone_number	HOD	Dept-ID

Professor

Professor_ID	Name	Phone-number

Course

Course_ID	Name	No.-of-credits	Dept-ID

Project

Project-ID	start_date	end-date	sponsorer	amount

Section

course_ID	section_ID	Timing	Class Room	Professor-ID

Enrollment

Roll-number	Section-ID

Project_Undertaken

Professor-ID	Project-ID

Prerequisite

Course_ID	Course-ID

<u>ex:</u>	<u>F</u>	<u>G</u>
	$B \rightarrow CD$	$B \rightarrow CDE$
	$AD \rightarrow E$	$B \rightarrow ABC$
	$B \rightarrow A$	$AD \rightarrow E$

<u>Ans:</u>	<u>F</u>	<u>G</u>	<u>F</u>	<u>G</u>
	$B^t = BCDEA$	$B^t = BACD$	$B^t = ACDB$	$B^t = ABCDE$
	$AD^t = ADE$	$= ABCDE$	$= ABCDE$	
		$AD^t = ADE$	$AD^t = ADE$	
				$'F'$ is covered by ' G '.
				So, ' F ' is equivalent to ' G '.

Q) How to identify key of a table from a closure set?

Identify key:

e.g:	$A \rightarrow BC$	$D \rightarrow BD$
	$CD \rightarrow E$	$ABH \rightarrow BC$
	$E \rightarrow C$	$DH \rightarrow BC$
		$D \rightarrow AEHT$

NOTE: In L.H.S., for individual attribute, we try to find closure set. (e.g. in this case, A, D, E)

$$A^t = ABC$$

$$B^t = ABDEH$$

$$E^t = EC$$

Out of the three, D^t contains all the attributes. Therefore, we identify 'D' as the key as it contains all the attributes.

Q) Find key for given functional dependencies: (try to find all combinations of attributes)

$A \rightarrow B$	$BCT = BCE$	$AB^t = AB$	$ACT = ACB \approx ABCE$
$BC \rightarrow E$		$BD^t = BD$	
$ED \rightarrow A$	$ED^t = AED = ABED$		

combine 3 combinations:

$$ACD^t = ACDB = ABCDE \quad (\text{usually we go for } ACD^t \text{ as single attribute of present})$$

$$BCD^t = BCDE = BCDEA = ABCDE$$

$$ECD^t = ABCDE$$

$A \rightarrow DE$	$A^t = ADE$	$AB^t = ABC$
$AB \rightarrow C$	$= ADEIJ$	$= ABCDE$
$B \rightarrow F$	$B^t = BF = BFGH$	$= ABCDEF$
$F \rightarrow GH$	$F^t = FGH$	$= ABCDFGH$
$D \rightarrow IJ$	$D^t = DJ$	$= ABCDEFGHIJ$

If we add AB^t with F or D or any other LHS attributes which give all attributes, we derive a set e.g.

candidate keys : $\{AB, ABF, ABD\}$



Super key (key along with other determinants)

Classification of Functional Dependencies:

(1) Partial F.D.:

→ Usually, non-key attribute depends upon key attributes. But where a particular non-key attribute depends on part of key, it is called a Partial F.D.

→ If any of the non-key attributes depends on part of the key, it is called as partial dependencies.

→ e.g.: $R = ABCD$ Key = AB
 $AB \rightarrow C$ $B \rightarrow D$ is a partial FD as 'D' is a nonkey & B is part of key.

$C \rightarrow D$
 $B \rightarrow D$ (B determines D or D depends on B)

→ Some situations where partial FD is not possible ..

→ e.g. If key is single attribute or Key contains all the attributes or if we have only two attributes present in table.

→ In the following condition, table can't have Partial F.D.:

- If the primary key consists of a single attribute.
- All the attributes in table are part of the key.
- If the table consists of only two attributes.

(2) Transitive F.D.:

→ If there is a relationship among non-key attributes, then they are called transitive FDs.

→ e.g.: $R = ABCD$
 $AB \rightarrow C$
 $C \rightarrow D$ (Transitive F.D. as {C,D,E} non-key set)
 $B \rightarrow D$
 AB is key.

→ In the following condition, table can't have transitive F.D.:

- All the attributes in table are part of the key.
- If table consists of only two attributes.

(3) Full Functional Dependency:

→ If there is a functional dependency $x \rightarrow y$, the removal of an attribute from x makes the F.D. $x \rightarrow y$ as invalid. Then we say that it is full F.D.

→ e.g.: $AB \rightarrow C$

remove B, so this is invalid.

$AB \rightarrow C$ is a full F.D.

IS. 08.2023

Q F: $AB \rightarrow CD$
 $C \rightarrow A$
 $D \rightarrow B$

$$AB^+ = ABCD \text{ (key)}$$

Check whether partial F.D. or transitive F.D. is there.

There is no partial F.D.

Transitive F.D. ($D \rightarrow B$)

If CD is the key, $C \rightarrow A$ & $D \rightarrow B$ are partial F.D.s

If partial F.D. or transitive F.D. are present, it is called anomalies in database.

Sol: split the table
(Till when?)

→ insertion/update/deletion anomaly.

Normalization:

It is a process of reducing redundancy from the universal database tables. Therefore, the insertion, deletion and update anomalies will be removed.

Types of Normal Form:

(1) First Normal Form:

If no redundant records exist in the table, then table is said to be in 1NF.
unique records

2. Second Normal Form (2NF)

If a table satisfies two criteria:

(a) It must be in 1NF

(b) There should be no partial dependency

It is said to be in second normal form.

$$\begin{array}{ll} \text{Example: } & AB \rightarrow C \\ \overline{\quad} & A \rightarrow DE \\ & B \rightarrow F \\ & F \rightarrow GH \\ & D \rightarrow IJ \end{array}$$

$$\begin{array}{l} AB \rightarrow C \\ AB^+ \rightarrow ABC \\ (\text{key}) = ABCDE \\ = ABCDEF \\ = ABCDEFGH \\ = ABCDEFGHIJ \end{array}$$

$$\begin{array}{l} A^+ = ADE \\ = ADEIJ \end{array}$$

$$\begin{array}{l} B^+ = BF \\ = BFGH \end{array}$$

(not in 2NF)

$A \rightarrow DE, B \rightarrow F$ has partial F.D.

* If there is a partial FD, we remove the partial dependent attribute from original table and place it in a separate table along with copy of its determinants.

R1: ABC (Determinant along with attributes for which A is key)

R2: ADEIJ (we are getting partial F.D.)

R3: BFGH (B is key)

Now we say this table is in 2NF.

$$\begin{array}{l} R1: A \rightarrow FC \\ C \rightarrow D \\ (B \rightarrow E) \\ \curvearrowright \text{partial F.D.} \end{array}$$

$$\begin{array}{l} R_1: AB \\ R_2: AFCD \\ R_3: BE \\ 2NF \end{array}$$

$$\begin{array}{l} AB^+ = AFCBE \\ = ABCDEF \\ (\text{key}) \end{array}$$

$$B^+ = BE$$

$$A^+ = AFCD$$

$$\begin{array}{l} R_1: AB \\ R_2: AFC \\ R_3: CD \\ R_4: BE \\ 3NF \end{array}$$

3) Third Normal Form (3NF)

If a table satisfies a criteria, it is said to be in 3NF:

(a) It should be in 2NF

(b) There should be no transitive F.D.

If there is a transitive dependency, we remove these transitive dependent attributes from the 2NF table and place it in a separate table along with copy of determinants.

3NF R₁: ABC

R₂: ADEIJ (A is key) \curvearrowright determines IJ
(transitive F.D.)

R₃: BFGH (B is key) \curvearrowright 'F' determines GH
(transitive F.D.)

3NF \curvearrowright

R₁: ABC

R₂: ADE

R₃: BF

R₄: DJ

R₅: FGH .. is in 3NF

4) Boyce Codd Normal Form (BCNF)

A table is said to be in BCNF if it is already in 3NF and all the determinants should be the key.

* 3NF can't handle the following situations:

- (i) If the table consists of multiple composite keys.
- (ii) If there is an overlapping attribute among candidate keys.

* Difference between 3NF and BCNF:

3NF

- (1) It focuses on primary key.
- (2) Still has redundancies.
- (3) If there is a dependency $x \rightarrow y$, it is allowed in 3NF, if either 'x' is a super key or 'y' is the part of key.
- Lossless decomposition: achieved

BCNF

- (1) It focuses on candidate keys.
- (2) More strict than 3NF
- (3) If there is a dependency $x \rightarrow y$, then it is allowed in BCNF if 'x' is a super key only.
- .. is hard to achieve.

Example of BCNF:

R(A-B-C-D), F.D.S: $A \rightarrow B$, $A \rightarrow C$, $C \rightarrow D$, $C \rightarrow A$.

Candidate keys: A and C

$$A^F = A B C D$$

$$C^F = A B C D$$

Transaction Processing:

Assumptions

Database server has single processor system.

2-tier DB architecture

Single DB system

→ While being carried out, transactions satisfy the ACID property:

- Atomicity
- Consistency
- Isolation
- Durability

• In application program usually contains number of simultaneous transactions.

• In user's point of view, transaction is a logical operation. It actually contains a set of micro-operations.

Atomicity: A transaction should either be done entirely or not at all performed.

• 'Recovery module' is responsible for enforcing atomicity. (By checking log files).

• Consistency: Application programmer should write efficient program that moves the database from one consistent to another consistent state, after the transaction has been performed.

• 'Application developer' is responsible for enforcing consistency.

• Isolation: • A transaction starts executing only after any previously executing transaction is completed. e.g. $A = 500$

• Isolation is enforced by 'concurrency control module'.

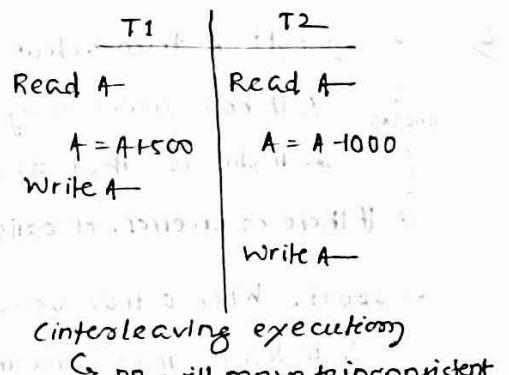
• A particular transaction cannot communicate with another.

• Isolation means separating the transactions.

• Durability: Enforced by 'Recovery Module'.

• The transactions that are going to be permanently recorded in the database must be successful.

• Log should be maintained only for fail cases and they must not be recorded in the database. \rightarrow DB will move to inconsistent state.



→ 1. (02.09.2023)

• Isolation operation:

→ If transactions are sequentially executed, their throughput will be low.

→ Certain regulated interleaf operations are allowed. (Ensured by concurrency control module)

→ In what scenarios database transaction fails?

↳ System crash (power failure, network issues)

↳ Transaction Program errors

↳ Transaction aborted by concurrency control manager, due to rule violation.

Partial Transaction Effects: Aborted

Successful Transaction: Recorded

→ All the details recorded in System log.

System log:

- Details of the running transactions are recorded in the log.
- Important resource for recovering from crashes.
- Always maintained on reliable secondary storage.

Log entries:

1. Beginning of transaction
2. Update operation, details - old values - new values
3. Ending transaction exit database

→ We are only concerned about read & write operations. Because the rest of the operations are handled by the application program.
 ↳ high level program

- From the database system point of view only the read and write operation of a transaction are important.
- Other operations happen in memory and do not affect the database on disk.

Notations:

- $R(x)$ - Transaction reads item ' x ' database object
- $W(x)$ - Transaction writes item ' x '

We can denote database objects as x, Y, Z , etc..

→ Commit: If a transaction issues a 'commit' command when

✓ checks 1. It has successfully completed its sequence of operations.

✓ 2. Indicates that its effect can be permanently recorded in the database.

③ If there is an error, it calls the concurrency control module to handle rule violations.

→ Abort: When a transaction issues 'abort.'

1. indicates that some internal error is there.

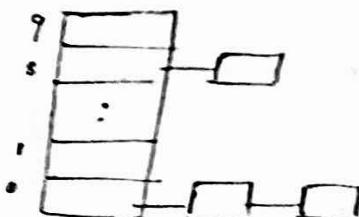
2. transaction wants to terminate in the middle.

• System obligations responsibilities when abort command has been issued:

→ Ensures that the partial work done by the transaction has no effect on the database on disk.

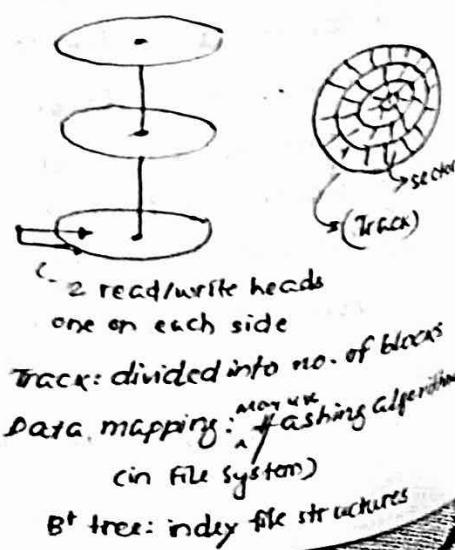
→ Database model for transaction processing:

- Granularity of item in database is dependent on the algorithm used. (as block size depends on algorithm)



(Use of Hashing)

Problem:
 Collision
 Solved by probing



- Transactions operate by exchanging the data with database only. They do not exchange messages between them.
- DB Model focuses on read/write/commit/abort operations and ignore the in memory operations.
- Transactions are not nested.

$R_i(x)$ i - stands for transaction number

→ read operation/transaction ' i ' on database item ' x '.

1. Disk block having ' x ' - copied to buffer page if required.
2. Required value(s) is assigned to program variable ' x '.

$DA = \text{Salary}$

*
memory variable → copied to buffer

$W_i(x)$ → write operation/transaction ' i ' on database item ' x '.

1. Disk block having ' x ' - copied to the buffer page if required.
2. Update the buffer value using program variable ' x '.
3. Transfer the block to the disk. can do immediately or later
if further modification is reqd.

C_i . Commit transaction ' i '

A_i - Abort transaction ' i '

→ 1- (07-09-2023) (11:00-12:00 pm)

- Transaction manager processes no. of transactions at a time to increase throughput and performance. Therefore in some cases we allow interleaved operations.
- There may arise anomalies, which need to be considered.
 - e.g.: Write operation on same database by two different simultaneous transactions.

❖ There are three anomalies:

- (1) RW anomaly
- (2) WR anomaly
- (3) WW anomaly

❖ Write-Read Anomaly:

<u>T1</u>	<u>T2</u>
1. $W(x)$	
	2. $R(x)$
	3. $x = x + 5;$
4. Revert old value of 'x'	
	5. $R(x)$

Interleaved operation being allowed, T1 & T2 are being executed. Though T2 updates x , as T1 reverts the old value, the second time T2 tries to read will obtain old value of ' x ' instead of getting updated value.

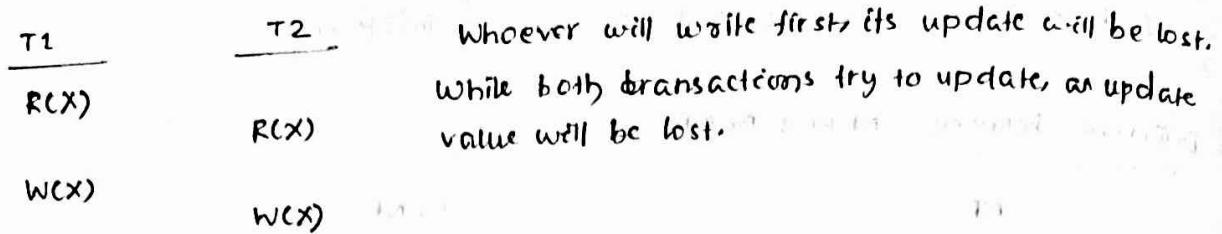
↳ Dirty Read operation.

❖ Read-Write Anomaly:

<u>T1</u>	<u>T2</u>
$R(x)$	
	$R(x)$
	$W(x)$
$R(x)$	

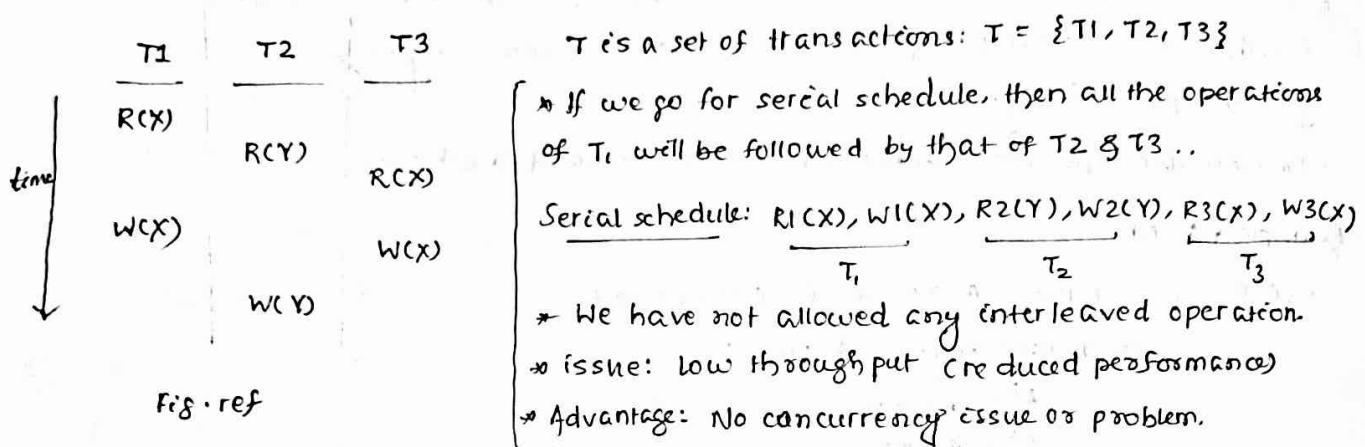
T1 is reading the database object ' x '. As T2 updates the value, the second time T1 tries to read will get a changed value. (Although T1 has performed no write or update operation)

* Write-Write Anomaly: (Lost update problem)



* Transaction Schedule: Interleaved operations will be allowed in such a way that no anomaly takes place.

→ It's a sequence of interleaved operations of a transaction set.



④ Serializability: Allow interleaving of operation maintaining some standard.

↳ Serializable schedule

- Interleaving of operations allowed.
- Equivalent to the serial schedules in some sense.

→ The effect of interleaving is same as that of some serial schedules.
i.e., interleaving equivalent to some serial schedules are allowed.

⑤ Conflicting pair of operations:

- Operations are involved on the same database item.
- These operations must belong to two different transactions.
- One of the operations must be the write operation.

Example: (Consider Fig.ref)

$R_1(X)$ and $W_3(X)$ involved in same DB item X

• belong to different transactions.

• one is write operation.

$R_2(X)$

.. $\{R_1(X), W_3(X)\}$

Other two conflicting pairs: $W_1(X)$ and $R_3(X)$

.. $\{R_3(X), W_1(X)\}$

$W_1(X)$ and $W_3(X)$

.. $\{W_1(X), W_3(X)\}$

• conflict serializable

• view serializable

[★ Design of healthcare management system
for chronic disease in Odisha]

Conflict Equivalence:

If schedule s_1 is said to be conflict equivalent to s_2 (denoted as $s_1 \equiv s_2$) if the relative order of any conflicting pair of operations is same in both s_1 and s_2 .

Consider example (fig ref):

$s_1: R_1(X), R_2(Y), \overset{R_3(X)}{\dots}, W_1(X), W_3(X), W_2(Y)$

$s_2: R_1(X), R_2(Y), R_3(X), W_2(Y), \overset{W_1(X)}{\dots}, W_3(X)$

Conflicting pairs:

$\{R_1(X), W_3(X)\}$

$\{R_3(X), W_1(X)\}$

$\{W_1(X), W_3(X)\}$

Relative orders of all conflicting pairs of operations in s_1 and s_2 are the same.

$\Rightarrow s_1$ is conflict-equivalent to scheduler s_2
or $s_1 \equiv s_2$.

$s_3: R_1(X), R_2(Y), W_1(X), R_3(X), W_3(X), W_2(Y)$

$(s_1, s_3): R_1(X), W_3(X)$ follow same relative order.

$R_3(X), W_1(X)$ have different relative order. - Out-of-pair violation
 $s_1 \not\equiv s_3$ (Though $w_1(X), w_3(X)$ have same relative order.)

(not conflict equivalent)

Conflict Serializable:

A schedule is called conflict serializable if it is conflict equivalent to some serial schedules.

(OR) If we can move non-conflicting operations such that

- The relative orders of operations are intact and the schedule becomes a serial schedule.

Example:

s1			s2		
T1	T2	T3	T1	T2	T3
R(X)			R(X)		
	R(Y)			R(Y)	
		R(X)			
W(X)			W(X)		
		W(X)		R(X)	
				W(X)	
W(Y)			W(Y)		

Conflicting pair of operations:

$\{R_1(X), W_3(X)\}$

$\{W_1(X), W_3(X)\}$

$\{R_3(X), W_1(X)\}$

have different

relative order

These are partly intact

For s_2 ,

$\{R_1(X), W_3(X)\}$

$\{W_1(X), W_3(X)\}$

$\{W_1(X), R_3(X)\}$

Transactions considered as node.
Arcs: connecting conflicting transactions.

* Build precedence graph.

* Topological sort.

* Result is the serial schedule.

Therefore s_2 is conflict serializable.

* If the graph is cyclic, non-conflict serializable.

If acyclic - is conflict serializable.

↳ View Serializability:

S_1 and S_2 are two schedules.

→ If schedule S_1 is view equivalent to schedule S_2 if

(i) T_i reads X in S_1 and also in S_2

(ii) T_i reads X written by T_j both in S_1 and S_2 .

(iii) Final write operation done by same transaction both in S_1 and S_2 .

example?

→ A schedule S_1 is said to be view serializable if it is view equivalent to a serial schedule S_2 .

↳ Concurrency control using locks:

→ If transaction requests for a lock on database item 'X' for both read and write operations. Then the transaction unlocks the database item 'X' once its job is over.

locks: binary locks (L/U)

→ At a particular time, on the same database item, a particular transaction will hold the locks:

transactions: $\{T_1, T_2\}$ database item: $\{X\}$

$T_1 \text{ L}(X)$

or

$T_2 \text{ L}(X)$

→ Lock/Unlock situations will be maintained by a 'locking schedule'.

↳ Locking and Serializability:

<u>T_1</u>	<u>T_2</u>	<u>T_3</u>	Lock Unlock Read Write
$L(X) R(X) U(X)$			
$L(X) W(X) U(X)$	$L(Y) R(Y) U(Y)$	$L(X) R(X) U(X)$	
		$L(X) W(X) U(X)$	
	$L(Y) W(Y) U(Y)$		

- Locking alone doesn't guarantee serializability.

- To attain serializability, find conflict equivalent or view equivalent.

↳ Two Phase Locking Protocol: (2 PL Protocol)

(Generate conflict equivalent schedule)

→ All lock requests of transaction precede with unlock request.

COR) A transaction has a locking phase followed by an unlocking phase.

→ If all the transactions follow the 2 PL protocol, the resulting schedule is conflict serializable.

Adv → Good for system performance. (Result is good)

Disadv → Deadlock may arise. ... use deadlock prevention/avoidance mechanism.

→ To detect the deadlock situation a wait-for graph is used.
.. detailed in OS.
(directed)

* Log Record Entries: (Details of the transaction)

- ↳ will be stored in reliable secondary storage.
- Who has updated? Timestamp? - Though we are not maintaining these two columns, log records keep them updated.

<start T> • Transaction T has started.

<End T> • Transaction T has completed.

<commit T> • Transaction has successfully completed the work
↳ changes must appear in database-disk.

<Abort T> • Transaction is failed

↳ changes should not be there in database division.

Update records • It is very specific to the logging methods.

Flush log • (Forcefully) Force-write the log entries onto the disk.

* Undo logging:

↳ database object
<T, X, V> → value

Transaction ↳ takes old value

→ Transaction 'T' has changed database item 'X' with old value 'V'.

* Rules: (Undo logging rules)

U1: <T, X, V>

If a transaction 'T' modifies the database item 'X', then the log record <T, X, V> must be written to the disk before the new value of 'X' is written to the disk.

U2: If a transaction 'T' commits, then the commit log record must be written to the disk only after all the database items changed by the transaction 'T' have been written to the disk.

For each modified database item X

{
send update entry <T, X, V> to disk
Write item X to disk}

}
Write <commit T> to disk

④ —
JAVA solves
surprise quiz

→ L- (09-09-2023) (09:00 - 10:00)

, Transaction T is doing money transfer of Rs. 100/- from account A to account B.

• Consistency requirement: A+B is same before and after the transaction.

m: memory variable

A: database item

Step	Action	M	Mem-A	Mem-B	Disk-A	Disk-B	Log
1					500	1500	
2	Read(A,m)	500	500		500	1500	
3	$m := m - 100$	400	500		500	1500	$\langle T, A, 500 \rangle$
4	Write(A,m)	400	400		500	1500	
5.	Read(B,m)	1500	400	1500	500	1500	
6.	$m := m + 100$	1600	400	1500	500	1500	$\langle T, B, 1500 \rangle$
7	Write(B,m)	1600	400	1600	500	1500	
8	Flush log				400	1500	
9	Output(A)	1600	400	1600	400	1600	
10	Output(B)	1600	400	1600	400	1600	
11							$\langle \text{commit}, T \rangle$
12	Flush log						

* Recovery using Undo Log:

Examine log and partition transaction into

→ Committed set : Transactions for which $\langle \text{commit}, T \rangle$ exists.

→ Incomplete set : Transactions for which $\langle \text{Abort}, T \rangle$ exists or $\langle \text{commit}, T \rangle$ doesn't exist.

Examine log in reverse direction:

For every update record $\langle T, Y, V \rangle$

T is committed: do nothing. All is well due to U_2

T is incomplete: Restore value of X on disk as V

- T might have changed some items on disk

- But log entries with old values are on disk due to U_1

Do $\langle \text{Abort}, T \rangle$ log entry for each incomplete T & flush log.

* Undo Logging: Crash Recovery:

Crash occurs sometimes before first Flush log.

DB on disk - unchanged.

T - recognized as incomplete.

Log entries of T: if present - used for "undo" $\langle \text{Abort}, T \rangle$

T - resubmitted.

If crash occurs after 11, $\langle \text{commit}, T \rangle$ - on disk : transaction completed
(not on disk : undoing performed).

If crash occurs after 12, NO action required.

* Redo Logging:

Undo logging: Cancels the effect of incomplete transactions

Ignores the committed transactions

All DB items to be sent to the disk before Txn can commit.

Results in lot of I/O operations, called **FORCE-waiting**.

No cannot be bulk.

Redo Logging :- Ignores incomplete Txns

- Repeats the work of committed Txns
- update log entry $\langle T, X, V \rangle$: v is the new value.

Buffer utilization might be down

Write-Ahead-Logging Rule

- Flow
- Update log records
 - commit records
 - changed DB items

8 : commit
9 : Flush log.

$\langle T, A, 400 \rangle, \langle T, B, 1600 \rangle$

Step	<u>Mem-A</u>	<u>Mem-B</u>	<u>Disk-A</u>	<u>Disk-B</u>	<u>log</u>
1					
2. Read(A,m)	500	500	500	1500	
3. m := m - 100	400	500	500	1500	
4. Write(A,m)	400	400	500	1500	$\langle T, A, 400 \rangle$
5. Read(B,m)	1500	400	500	1500	
6. m := m + 100	1600	400	500	1500	
7. Write(B,m)	1600	400	1600	500	$\langle T, B, 1600 \rangle$
8.					
9. Flush log					
10. Output(A)	1600	400	400	1500	
11. Output(B)	1600	400	1600	400	

+ Recovery using Redo log:

→ Examine logs & partition Txns into Committed set and Incomplete set.

→ Examine log in forward direction:

for every update record $\langle T, X, V \rangle$

T is Incomplete: Do nothing. DB on disk has no effects.

T is Committed: Unsure if all the effects of T are reflected on Disk

- But log entries with new values are on disk (WAL)

- Redo the changes as per the log entry.

Do Abort T for each incomplete T and Flush log.

- Crash Recovery:

before step-8 : $\langle \text{commit } T \rangle$ not made by T : recognized as incomplete: No actions required.
DB on disk - not changed. (WAL) (Abort T, resubmit.)

after step-8 : $\langle \text{commit } T \rangle$ on disk - T is redone with the help of log entries

(not on disk - incomplete - do nothing - WAL - Abort T & resubmit T.)

- Undo-Redo logging:

Better flexibility, uses more detailed logging.

$\langle T, X, U, V \rangle$

$\begin{matrix} \text{old} \\ \text{value} \end{matrix}$ $\begin{matrix} \text{new} \\ \text{value} \end{matrix}$

• Before modifying any DB item X on disk, because of changes made by some txn T, it is necessary that the transaction log $\langle T, X, U, V \rangle$ must appear on DB.

- $\langle \text{Commit } T \rangle$ & disk changes - in any order

8: Flush log
 9: Output log
 10: Commit log
 11: Output R

* Recovery using Undo-Redo logs:

Redo all committed Txns - order - earliest first
 Undo all incomplete Txns - order - latest first

both necessary

* Crash Recovery:

before $\langle \text{commit}, T \rangle$ incomplete T - undone
 after $\langle \text{commit}, T \rangle$ committed T - redone

* Recoverable Schedule:

Serializability and Recoverability are orthogonal concepts.

\hookrightarrow no restriction on commit
 \hookrightarrow no restriction on locking

- Cascadeless
- cascading

Some schedules have cascading effect

Lab: Txn management in SQL

Next: Tuple Relational Algebra

\rightarrow L- (15.09.2023) (11:00-12:00)

* Relational Algebra:

1) Select (σ)
 Project (π)
 Division (\div)
 Rename (ρ)
 Assignment (\leftrightarrow)
 Join (\bowtie)

Basic Relational Algebra functions

\rightarrow SQL engine converts SQL statements to Relational Algebra form

- 2) Union (\cup), Intersection (\cap), Difference (\setminus) and Cartesian Product (\times)
 3) Extended relational algebraic functions:
 Minimum, Maximum, Average, Count

\hookrightarrow Select (σ)

↳ Unary operator

↳ Syntax: $\sigma_{\text{cond}}^{\text{rel}}$

$\sigma_{\text{cond}}^{\text{rel}}$ \hookrightarrow Relation

↳ e.g.: Suppose we want to select students of branch computer science & Engineering with age greater than 19.

$\sigma_{\text{branch} = "CSE" \text{ and } \text{age} > 19}^{\text{student}}$ - will filter the rows only of a relation

- \rightarrow The conditions are (commutative) allowed to be interchanged (order doesn't affect result)
 \rightarrow It is commutative in nature.

$$\sigma_{\langle c_1 \rangle \text{ and } \langle c_2 \rangle}^{\text{rel}} = \sigma_{\langle c_2 \rangle \text{ and } \langle c_1 \rangle}^{\text{rel}}$$

$$\begin{aligned} \rightarrow \delta_{\langle c_1 \rangle} (\delta_{\langle c_2 \rangle} (R)) &\equiv \delta_{\langle c_2 \rangle} (\delta_{\langle c_1 \rangle} (R)) \\ &\equiv \delta_{\langle c_2 \rangle \text{ and } \langle c_1 \rangle} (R) \end{aligned}$$

conjunction

composite condition: AND, OR, NOT

\rightarrow The comparison operators $<, \leq, >, \geq, =, \neq$ can be used in the condition clause.

• Project (π)

↳ Unary operator

↳ Syntax: $\pi_{\langle \text{attributes} \rangle} (R)$

↳ Filters the columns

\rightarrow e.g. $\pi_{\text{name}, \text{age}} (\text{student})$

\rightarrow Commutative is not allowed in project statement.

If allowed some properties of table will be lost.

name	age
:	:

all rows

SQL

SELECT name, age
FROM students
WHERE branch = "CSE"
AND age > 19;

convert into Relational Algebra using select & project

$\pi_{\text{name}, \text{age}} (\delta_{\text{branch} = "CSE" \text{ and } \text{age} > 19} (\text{students}))$

Nested Query should show the desired output.

- If all of the operations are performed on the top of table, then order won't matter usually.
- ② Go in such a way that inner query should produce sufficient data for outer query.

• Division (\div)

↳ binary operator

$$R = ABC$$

$$S = AB$$

$$R \div S = C$$

↳ Subset of R-S

eg:

R			S		R \div S	
A	B	C	C	A	B	C
a ₁	b ₁	c ₁	c ₁	a ₁	b ₁	c ₁
a ₂	b ₁	c ₂	c ₁	a ₂	b ₁	
a ₂	b ₂	c ₁	c ₁	a ₂	b ₂	
a ₁	b ₂	c ₂	c ₂	a ₁	b ₂	
a ₂	b ₁	c ₃	c ₃	a ₂	b ₁	
a ₁	b ₂	c ₄	c ₄	a ₁	b ₂	
a ₁	b ₁	c ₅	c ₅	a ₁	b ₁	

eg. 2

S		R				R \div S	
A	B	A	B	C	D	C	D
a ₁	b ₁	a ₁	b ₁	c ₁	d ₁	a ₁	d ₁
a ₂	b ₂	a ₂	b ₂	c ₁	d ₁		
a ₁	b ₁	a ₁	b ₂	c ₂	d ₂	c ₂	d ₂
a ₁	b ₂	a ₁	b ₁	c ₃	d ₃	c ₃	d ₃
a ₂	b ₂	a ₂	b ₁	c ₃	d ₃		

Ans:

C	D
a ₁	d ₁
a ₂	d ₂
c ₃	d ₃

Duplicates must not be kept

eg. 3

↳ Union, Intersection, Difference:

R		
A	B	C
a ₁	b ₁	c ₁
a ₂	b ₂	c ₃
a ₂	b ₁	c ₂

S		
A	B	C
a ₁	b ₁	c ₁
a ₁	b ₁	c ₂
a ₂	b ₂	c ₃
a ₃	b ₂	c ₃

RUS		
A	B	C
a ₁	b ₁	c ₁
a ₁	b ₂	c ₃
a ₂	b ₁	c ₂
a ₁	b ₁	c ₂
a ₃	b ₂	c ₃

RNS		
A	B	C
a ₁	b ₁	c ₁
a ₁	b ₂	c ₃

R - S

A	B	C
a ₂	b ₁	c ₂

S - R

A	B	C
a ₁	b ₁	c ₂
a ₃	b ₂	c ₃

↳ Intersection (\cap) can be expressed using Union and Difference:

$$R \cap S = (R \cup S) - (R - S) \cup (S - R)$$

↳ Join (\bowtie) can be expressed using select and cartesian product:

$$R \bowtie_c S = \sigma_{c \in C} (R \times S)$$

↳ Division (\div) can be expressed using project, product and Difference:

$$R \div S = T$$

$$T_1 \leftarrow \pi_{\{A\}} R$$

$$T_2 \leftarrow \pi_{\{A\}} (S \times T_1) - R$$

$$T = T_1 - T_2$$

↳ Basic operators are:

a. Select (σ)
b. Project (π)

c. Union (\cup)

d. Difference (\setminus)

e. Cross Product (\times)

Next class:

Query Optimization /

Q1

T ₁	T ₂
R(A)	
W(A)	R(B)
	W(B)
R(B), W(B)	R(A) W(A)

Schedule S is:

- (i) view serializable but not conflict serializable
- (ii) conflict serializable but not view serializable
- (iii) Both view & conflict serializable
- (iv) Neither view nor conflict serializable

T ₁	T ₂
R(O)	
W(O)	W(O)

Q2 Options same as those of Q1

Read Transaction management
G Test tomorrow