

Name	Product_name	Proj_name	Category	Supplier	Customer
smith	X (01)	Proj 1	Electronics	Supplier A	Customer 1
smith	X (01)	Proj 2	Electronics	Supplier B	Customer 2
Adamsky	X (01)	Proj 3	Electronics	Supplier C	Customer 3
Wallon	S (01)	Proj 4	Electronics	Supplier D	Customer 4
Adamsky	X (01)	Proj 5	Electronics	Supplier E	Customer 5
	X (01)	Proj 6	Electronics	Supplier F	Customer 6

The dependency relation is  $R_1 \Delta R_2 \Delta R_3 = \text{Supply}$  and we don't have any non-trivial join dependency in  $R_1, R_2 and  $R_3$ .$

#### MODULE - IV

DT-301 09/24

Query processing refers to the range of activities that are involved in extracting data from a database.

#### Basic steps in Query processing :-

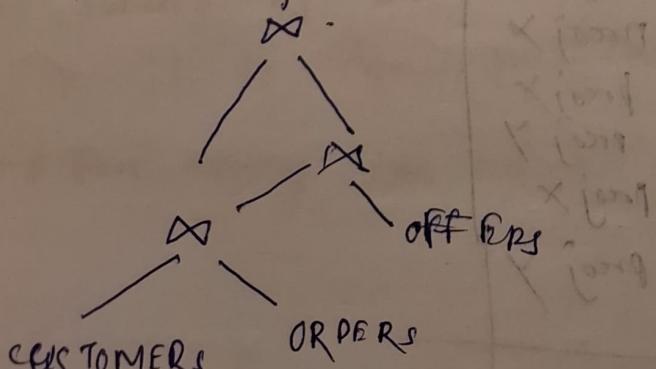
The basic steps involved in processing query are:-

- 1) Parsing and Translation.
- 2) Optimization

- 3) Evaluation

#### 1) Parsing and Translation:-

TT.   
 Price > 5000.



Information about relations and attributes:-

$N_R$  :- No. of tuples in relation R

$B_R$  :- No. of blocks that contain tuples of relation R.  
(where the data is stored)

$S_R$  :- size of tuple R

$F_R$  :- blocking factor; no. of tuples from R that fit  
into one block ( $F_R = \left\lceil \frac{N_R}{B_R} \right\rceil$ )  
Scaling  $\rightarrow b_lg - 0$

$V(A, R)$  :- no. of distinct values for attribute A in  
relation R.

$SCC(A, R)$  :- selectivity of attribute A

2 Avg no. of tuples of R that satisfy an  
equality condition on A

$$SCC(A, R) = \frac{t_R}{V(A, R)}$$

Note: The size of tuple means the amount of  
memory (in bytes) required for each tuple.

Measurement of query cost:-

Number of disk seeks no. of blocks read,  
Number of blocks written.  
Seek time is the amount of time that  
a disk drive head take to move to a specific  
location of disk.

Disk seek is the actual physical positioning of  
the re-drive head of the disk.

Disk transfer occurs from main memory to secondary memory  
suppose

$t_F \rightarrow$  time for transferring one block

$t_S \rightarrow$  time for one disk seek operation

cost for  $b$  block transfers +  $S$  seek operations

$$= b \times t_F + (S \times t_S)$$

### Selection operation

$S_1 \rightarrow$  linear search  $\rightarrow$  for both sorted and unsorted list

$S_2 \rightarrow$  binary search  $\rightarrow$  only for sorted list

$\downarrow$  words on divide and conquer method.

It is much more faster than linear search.

No of blocks that needs to be examined to find a block containing the required record is  $\lceil \log_2(B_R) \rceil$

$B_R \rightarrow$  No of blocks.

cost of binary search, i.e., the file ordered by

~~for~~ If the selection condition is base on non

~~key~~ attributes the cost ~~for  $S_2$~~

$$\text{cost}(S_2) = \lceil \log_2(B_R) \rceil + \left\lceil \frac{\text{sc}(A, R)}{F_R} \right\rceil - 1$$

no of blocks required for storing the result of selection operation

For key attribute:

$$\text{cost}(S_2) = \lceil \log_2(B_R) \rceil$$

then  $\text{sc}(A, R) = 1$

If primary then  $\text{sc}(A, R) = 1$

We consider a employee relation where  
10,000 tuples. If blocking factor,  $F_R = 10$   
employee

VC deptno, employee  $\leq 50$

↳ The no. of distinct department existing  
on the employee database.

Assume the following rel<sup>n</sup> operation.

$O_{deptno=20}$  (employee) and employee is sorted  
on such key deptno. Find the cost of selecting

the required records that satisfy selection  
criteria using Linear search and binary search

Linear search

$$cost(S_1) = \log_2 \frac{N_{Rmp}}{2} + \log_2 \frac{N_{Rmp}}{2} + \log_2 \frac{N_{Rmp}}{2}$$

$$B_{Rmp}^2 = \frac{N_{Rmp}}{F_{Rmp}} = \frac{10000}{10} = 1000$$

$$\text{Linear search} = S_1 = 1000$$

Binary search

$$\begin{aligned} cost(S_2) &= \log_2 (B_R) + \frac{SC(CAR)}{FR} - 1 \\ &= \log_2 1000 + \frac{200}{10} - 1 = 3 + 20 - 1 \\ &= 22 \end{aligned}$$

$$SC(CAR) = \frac{NR}{V(CAR)} = \frac{10000}{50} = 200$$

Q) The emp rel<sup>n</sup> is sorted on search key  
dept no.

There is an equal no. of employees present in  
each department.

## Selection operation using Indices :-

[DT-22/24]

Index scan :-

### NOTATION

$HT_I$ : number of levels in index I ( $B^+$ -tree)

$LB_I$ : number of blocks occupied by leaf node  
in index I (first-level blocks)

$Nval_I$ : number of distinct values for the search key.

$B^+$  tree is balanced binary search tree.

It follows multi-level index format.

### S3

$S3$  is a selection operation using index scan.

$S3$  - primary index for A, A primary key, with  
equality comparison  $A = a$  (as the selection criteria)

Ex:- select \* from student where std-id = 105

$$\text{cost}(S3) = HT_I + 1 \quad (\text{only 1 tuple satisfies the condition})$$

$HT_I \rightarrow$  no of levels in the  $B^+$  tree that

is used to store the index I.

### S4

$S4$  - primary index for A, A non-key attribute  
with equality comparison  $A = a$  (as the selection criteria)

$$\text{cost}(S4) = HT_I + \left[ \frac{\text{sc}(A, R)}{F_R} \right]$$

### S5

$S5$  - non-primary/secondary index on non-key  
attribute A with equality comparison  $A = a$

$$\text{cost}(S5) = HT_I + SC(A/R)$$

selection involving comparisons

$S_6$  - selection involving a primary index on A

$S_7$  - selection involving a secondary index on A

$S_6$

Ex: select \* from student where std-fd  $\geq 105$

$$\text{cost}(S6) = \text{seek } HT_I + \text{B. number of block before } ( \leq, \leq ) \text{ are after}$$

$( >, \geq )$  is added

$S_7$  (It is same as  $S_5$  but here operator is comparison operator)

Ex: select \* from student where age  $\geq 18$

$$\text{cost}(S7) = HT_I + SC(A/R) + \text{Number of blocks searched}$$

## Q) Note & employee

Ex:- Assume the query CUSTOMERS  $\bowtie$  ORDERS (with join attribute only being CName)

-  $N_{CUSTOMERS} = 5000$  tuples  
 $F_{CUSTOMERS} = 20$ , i.e  $B_{CUSTOMERS} = 5,000/20 = 250$  blocks

-  $N_{ORDERS} = 10,000$  tuples  
 $F_{ORDERS} = 25$  i.e  $B_{ORDERS} = 10,000 / 250 = 400$  blocks

-  $V(CName, ORDERS) = 2500$   
 The cartesian product  $R \times S$  requires  $N_R \times N_S$  tuples.  
 $s_R$  = size of individual tuple in  $rel^R$

each tuple requires  $s_R + s_S$  bytes

$s_S$  = size of individual tuple in  $rel^S$

size of join using information about foreign key

If schema(R)  $\bowtie$  schema(S) = foreign key in S referencing R, then the no. of tuples in  $R \bowtie S$  is exactly  $N_S$ .

→ In the example query CUSTOMER  $\times$  ORDERS, the result has exactly  $N_S = 10000$  tuples.

size of join without information about foreign key

If schema(R)  $\bowtie$  schema(S) = {A} which is not a key for R or S, then the no. of tuples:-

$$\min\left(\frac{N_R \times N_S}{V(A, S)}, \frac{N_R \times N_S}{V(A, R)}\right)$$

A = At is common in both R and S.

In previous ex:-

$$N_R = 5000, N_S = 10000$$

$$V(A, S) = 5000, V(A, R) = 2500$$

$$\min\left(\frac{5000 \times 10^4}{5000}, \frac{5000 \times 10^4}{2500}\right)$$

$$\min(10^4, 20^4) = \min(10000, 20000)$$

No. of tuples = 10000

Nested loop join:-

It is a join that contains a pair of nested loops.

representation:-

$R \bowtie_C S$

$C \rightarrow$  condition

$\rightarrow R$  is outer rel<sup>n</sup> and  $S$  is inner rel<sup>n</sup> of the join

cost Analysis of nested loop join :-

worst case :- The database buffer can hold only one block of each rel<sup>n</sup>

In this case the no. of block requires.

$\rightarrow BR + NR * BS$  disk accesses.

Best case :- The database buffer

there enough space for both rel<sup>n</sup> to fit into the database buffer simultaneously, ~~so no thrash~~.

$\rightarrow BR + BS$  disk accesses.

Block Nested loop join :-

~~Here~~ here the for each block no. of people are present.

Cost Analysis :- Worst case :-

Database buffer can hold only one block of each rel<sup>n</sup>

$\rightarrow BR + BR * BS$  disk access.

Best case :-

both rel<sup>n</sup>'s fit into database buffer

$\rightarrow BR + BS$  disk access.

## Ex-1 Nested loop join

worst case:-  $BR + NR * BS$

$$= 250 + 5000 \times 400$$

$$= 250 + 2000000$$

$$= 2000250$$

best case:-  $BR + BS$

$$= 250 + 400$$

$$= 650$$

## Block nested loop join :-

worst case:-  $BR + BR * BS$

$$= 250 + 250 \times 400$$

$$= 250 + 100000$$

$$= 100250$$

best case:-  $BR + BS$

$$= 250 + 400$$

$$= 650$$

## Index Nested loop join :-

If it is used to speed up the process of joining two tables.

### cost ANALYSIS:-

#### Worst case :-

Database space buffer has space only for one block of R and one block of index associated with S.

If cost of a single selection on S using the join condition. Then cost of join =  $BR + NR^c$

(Q) CUSTOMER ORDERS with CUSTOMER → our  
 Let orders have primary B+ tree index on the join attribute  
 CName let the height of the B+ tree to accommodate  
 is 4?

$$SC(A/s) = \frac{Ns}{B(A/s)} = \frac{10000}{2500} = 4$$

what will be c for search for & one access is  
 cost c = height of the tree + need to get the a

$$= 4 + 1 = 5$$

$$\begin{aligned} \text{cost} &= BR + NR * c \\ &= 250 + 5000 * 5 = 250 + 25000 \\ &= 25250 \end{aligned}$$

sort merge join:-

It is the combination of two processes i.e sorting and merging.

1) Step 1:- If both the relns are sorted.  
 base on the join attribute

cost analysis

If the reln's are sorted and tuples having the same value on the join attribute are placed consecutively then we need to read each tuple only once and thus the disk blocks are also read once.  
 Worst case :-

The no. of disk access =  $BR + B_s * B_s$

best case:-

The no. of disk access =  $BR + B_s + \text{cost of sorting}$

Q) CUSTOMER ORDERS with CUSTOMER  $\rightarrow$  outer reln  
 Let orders have primary BT tree index on the join attribute  
 CName let the height of the BT tree to accommodate  
 is 4?

$$sc(A/s) = \frac{Ns}{B(A/s)} = \frac{10000}{2500} = 4$$

what will be c to search page  $\leftarrow$  one more access is  
 cost c = weight of the tree + need to get the a  
 $= 4 + 1 = 5$

$$\begin{aligned} \text{cost} &= BR + Nrd * c \\ &= 250 + 5000 * 5 = 250 + 25000 \\ &= 25250 \end{aligned}$$

### sorted merge join:-

If is the combination of two processes i.e sorting and merging.

Step 1:- If both the reln's are sorted, base on the join attribute.

Step 2:- join

### cost analysis

If the reln's are sorted and tuples having the same value on the join attribute are placed consequitively then we need to read each tuple only once and thus the disk blocks are

also read once.  
worst case:-

The no. of disk access =  $BR + \frac{BR * Bs}{2}$

best case:-

The no. of disk access =  $BR + Bs + \text{cost of sorting}$

## Query optimization

Find the names of all customers who have an account at any branch located in Brooklyn.

$\Pi_{\text{customer-name}}(\sigma_{\text{branch-city} = \text{"Brooklyn"} }(\text{branch}))$

$\Pi_{\text{customer-name}}(\sigma_{\text{branch-city} = \text{Brooklyn}}(\text{branch}))$   
 (account  $\bowtie$  depositing)

### Notations:

$\phi_i$  → represents predicate

$L_i$  → used to denote the list of attributes

$E_i$  → denotes the relational-algebra expression.

$\pi$  → relation name

Rules  
 1) consecutive selection operations can be decomposed into a sequence of individual selections.

$$\sigma_{\phi_1 \wedge \phi_2}(E) = \sigma_{\phi_1}(\sigma_{\phi_2}(E))$$

2) selection operations are ~~completely~~ commutative

$$\sigma_{\phi_1}(\sigma_{\phi_2}(E)) = \sigma_{\phi_2}(\sigma_{\phi_1}(E))$$

3) only the final operation in a sequence of projection operations is required, the rest can be omitted

$$\Pi_{L_1}(\Pi_{L_2}(\dots \Pi_{L_n}(E) \dots)) = \Pi_{L_1}(E)$$

4) insertion operation can be combined with cartesian product and theta joins in the following manner

$$\sigma_Q(E_1 \bowtie Q E_2) = E_1 \bowtie Q E_2$$

$$\sigma_Q(E_1 \bowtie_{Q_1} E_2) = E_1 \bowtie_{Q_1} \sigma_{Q_2} E_2$$

5) Theta-join operations are commutative

$$E_1 \bowtie_Q E_2 = E_2 \bowtie_Q E_1$$

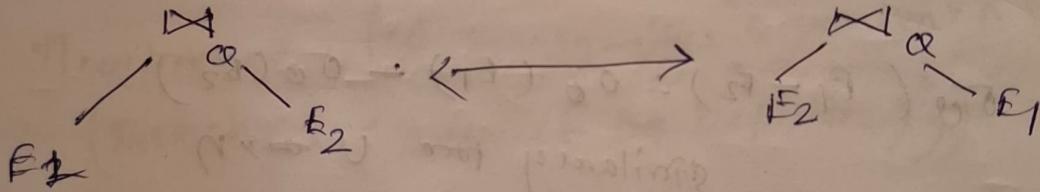
6) a) Natural join operations are associative

$$(E_1 \bowtie_{Q_2} E_2) \bowtie_{Q_3} E_3 = E_1 \bowtie_{Q_1} (E_2 \bowtie_{Q_3} E_3)$$

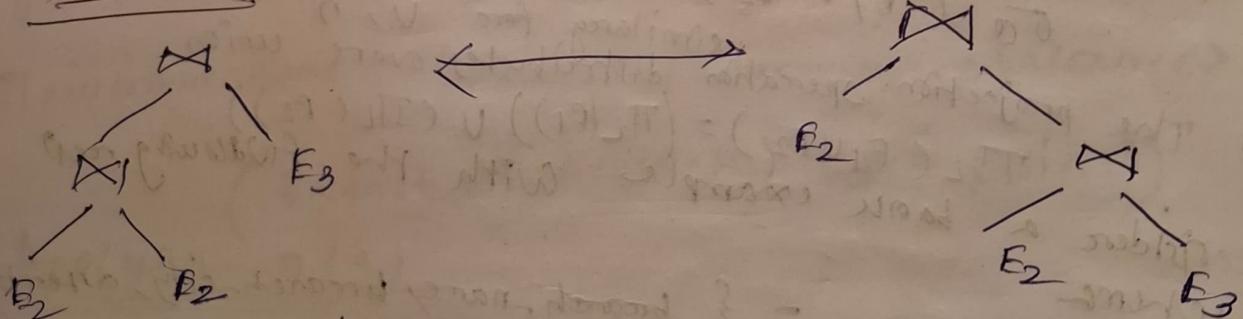
# b) Theta-join operations are associative

$$(E_1 \bowtie_{Q_1} E_2) \bowtie_{Q_2 \bowtie_{Q_3}} E_3 = E_1 \bowtie_{Q_1} (E_2 \bowtie_{Q_3} E_3)$$

Rule - 5



Rule - 6 (a)



a) When all the attributes in  $\alpha$  involve only the attributes of one expression.

b)  $\sigma_Q(E_1 \bowtie_{Q_1} E_2) = (\sigma_Q(E_1)) \bowtie_{Q_1} E_2$  when  $Q_1$  involves only the attributes of  $E_1$  and  $Q_2$  involves only the attribute of  $E_2$ .

b)  $\sigma_{Q_1 \bowtie_{Q_2} E_2}(E_1 \bowtie_Q E_2) = (\sigma_{Q_1}(E_1)) \bowtie_{Q_2} (\sigma_{Q_2}(E_2))$

(a) Let  $L_1$  and  $L_2$  be attributes of  $E_1$  and  $E_2$  respectively. If  $L_1$  involves only attributes from  $E_1$  and  $L_2$  involves only attributes from  $E_2$ .

(b)  $\pi_{L_1 \cup L_2}(E_1 \bowtie_Q E_2) = (\pi_{L_1}(E_1)) \bowtie_Q (\pi_{L_2}(E_2))$  from ppt

$$b) \pi_{L_1 U L_2}(E_1 \Delta E_2) = \pi_{L_1 U L_2}((\pi_{U_{L_2}}(E_1)) \Delta$$

$$\pi_{L_2 U L_1}(E_2))$$

q) The set operations union and intersection are commutative.

$$E_1 \cup E_2 = E_2 \cup E_1$$

$$E_1 \cap E_2 = E_2 \cap E_1$$

set difference is not commutative

10) set union and intersection are associative

$$(E_1 \cup E_2) \cup E_3 = E_1 \cup (E_2 \cup E_3)$$

$$(E_1 \cap E_2) \cap E_3 = E_1 \cap (E_2 \cap E_3)$$

ii) The selection operation distributes over  $\cup$ ,  $\cap$ , and -

$$\sigma_a(E_1 - E_2) = \sigma_a(E_1) - \sigma_a(E_2)$$

similarly for  $\cup$  or  $\cap$

$$\sigma_a(E_1 - E_2) = \sigma_a(E_1) - E_2$$

similarly for  $\cup$ ,  $\cap$

11) The projection operation distributes over union

$$\pi_L(E_1 \cup E_2) = (\pi_L(E_1)) \cup (\pi_L(E_2))$$

We consider a bank example with the following ref schema

Branch schema = { branch-name, branch-city, address }

Account schema = { account-number, branch-name, balance }

Depositor schema = { customer-name, account-number }

Find the names of all customers who have an account at any branch located in "Brooklyn".

Branch(AccOUNT Depositor)

$\Pi$  customer-name ( $\sigma_{\text{branch-city} = \text{"Brooklyn"}} (\text{Branch} \bowtie \text{Account} \bowtie \text{Depositor})$ )  
by applying rule 7(a) this query can be transformed  
to an equivalent query that generates smaller  
intermediate rel?

$\Pi$  customer-name ( $\sigma_{\text{branch-city} = \text{"Brooklyn"} \wedge \text{balance} > 1000} (\text{Branch} \bowtie \text{Account} \bowtie \text{Depositor})$ )  
suppose we modify the original query to restrict  
our search on customers who have a  
balance greater than \$1000 dollars

$\Pi$  customer-name ( $\sigma_{\text{branch-city} = \text{"Brooklyn"} \wedge \text{balance} > 1000} (\text{Branch} \bowtie \text{Account} \bowtie \text{Depositor})$ )  
(Branch  $\bowtie$  (Account  $\bowtie$  Depositor))

by applying rule 6(a)

$\Pi$  customer-name ( $\sigma_{\text{branch-city} = \text{"Brooklyn"} \wedge \text{balance} > 1000} (\text{Branch} \bowtie \text{Account} \bowtie \text{Depositor})$ )  
(Branch  $\bowtie$  Account  $\bowtie$  Depositor)

Applying 7(a)

$\Pi$  customer-name ( $\sigma_{\text{branch-city} = \text{"Brooklyn"} \wedge \text{balance} > 1000} (\text{Branch} \bowtie \text{Account} \bowtie \text{Depositor})$ )  
(Branch  $\bowtie$  Account  $\bowtie$  Depositor)

Applying rule 1

$\Pi$  customer-name ( $\sigma_{\text{branch-city} = \text{"Brooklyn"} \wedge \text{balance} > 1000} (\text{Branch} \bowtie \text{Account})$ )

### Schedule - 3

$$A = 1000$$

$$A = 1000 - 50 \\ = 950$$

$$B = 2000$$

$$B = 1980$$

$$200 \times \frac{1}{100}$$

$$950 \times \frac{1}{100}$$

$$A = 950$$

$$A = 1000 - 950 = 950 - 95$$

[DT:- 04/11/2021]

### Schedule - 4

$$1) A = A - 50$$

$$= 1000 - 50 = 950$$

$$A = 1000$$

$$A = 1000 - 100 \\ = 900$$

$$B = 2000$$

$$B = 2050$$

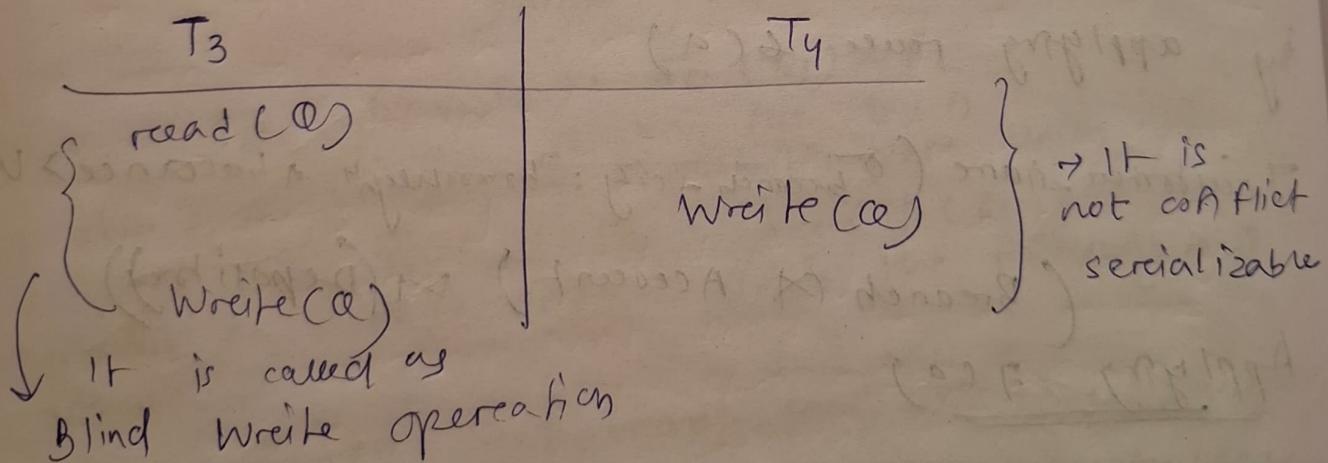
$$B = 2150$$

$$A + B = 3050$$

Conflict serializability

given

ex:- schedule - 1 & schedule - 2



View serializability :-

T <sub>3</sub>	T <sub>4</sub>	T <sub>6</sub>
read(Φ)	Write(α)	
Write(α)		write(α)

The corresponding serial schedule of this concurrent schedule.

→ There is no read<sub>2</sub> operation after write<sub>2</sub>  
That's why condition 2 is not applicable.

Schedule 9 is view equivalent to the serial schedule.

It is view equivalent schedule but not conflict serializable schedule.

The schedule containing blind writes is view serializable schedule that is not conflict serializable.

Schedule 9 and  $\{T_3, T_4, T_6\}$  are not conflict serializable because every pair of instruction are not conflict each other.

Testing for serializability:-

steps

- 1) Input is concurrent schedule S
- 2) Output is whether the schedule S is conflict serializable or not.

TAS  
Construct the precedence graph for schedule 1, schedule 2, schedule 3 and schedule 9.

Suppose initially  $A = \$100$  and  $B = \$200$

### Recovery system

- Failure classification
- 1) Transaction Failure :-  
a) logical failure  
b) system failure
- 2) system crash
- 3) Disk Failure

read(A)	read C
$A = A - 50$	$C = C - 100$
write(A)	write(C)
read(B)	
$B = B + 50$	
write(B)	

$$A = 100, \quad B = 200, \quad C = 700$$
$$\downarrow \qquad \downarrow \qquad \downarrow$$
$$A = 950 \quad B = 2050 \quad C = 650$$

### Checkpoints

- When a system failure occurs we must consult the log to determine those transaction that needs to be redone and those that need to be undone.
- We need to search the entire log to decide the recovery procedure.

There are 2 major difficulties with this approach

- (1) the search process is very time consuming
- (2) the transaction that needs to be redone have already written their updates into the database so redoing such transaction causes the recovery to take longer time, so to reduce this type of overhead check-points are introduced.

During execution the system maintains the log using defferred database modification or immediate database modification

In addition the system periodically performs check points

which requires the following sequence of steps :-

step 1 :- output onto stable storage all log records  
currently residing in main memory.

step 2 :- output to the disk all modified buffer  
blocks

step 3 :- output onto the stable storage a log  
recorded & checkpoint