



Module 17: Transactions

Database System Concepts, 7th Ed.

©Silberschatz, Korth and Sudarshan

See www.db-book.com for conditions on re-use



Outline

- Transaction Concept
- Transaction State
- Concurrent Executions
- Serializability
 - Conflict Serializability
 - View Serializability
- Test for Serializability



Transaction Concept

- A **transaction** is a *unit* of program execution that accesses and possibly updates various data items.
- E.g., transaction to transfer \$50 from account A to account B:
 1. **read**(A)
 2. $A := A - 50$
 3. **write**(A)
 4. **read**(B)
 5. $B := B + 50$
 6. **write**(B)
- Two main issues to deal with during transaction execution:
 - **Failures** of various kinds, such as hardware failures and system crashes
 - Concurrent execution of multiple transactions



ACID Properties of a Transaction

To preserve the integrity of data, the database system must maintain the following properties of transactions:

- **Atomicity (A):** Either all operations of the transaction are properly reflected in the database or none are.
- **Consistency (C):** Execution of a transaction in isolation preserves the consistency of the database.
- **Isolation (I):** Although multiple transactions may execute concurrently, each transaction must be unaware of other concurrently executing transactions. Intermediate transaction results must be hidden from other concurrently executed transactions.
 - That is, for every pair of transactions T_i and T_j , it appears to T_i that either T_j finished execution before T_i started, or T_j started execution after T_i finished.
- **Durability (D):** After a transaction completes successfully, the changes it has made to the database persist, even if there are system failures.



Example of Fund Transfer

- Transaction to transfer \$50 from account A to account B:
 1. **read**(A)
 2. $A := A - 50$
 3. **write**(A)
 4. **read**(B)
 5. $B := B + 50$
 6. **write**(B)
- **Atomicity requirement**
 - If the transaction fails after step 3 and before step 6, money will be “lost” leading to an inconsistent database state
 - Failure could be due to software or hardware failure.
 - The system should ensure that updates of a partially executed transaction are not reflected in the database
- **Durability requirement** — once the user has been notified that the transaction has completed (i.e., the transfer of the \$50 has taken place), the updates to the database by the transaction must persist even if there are software or hardware failures.



Example of Fund Transfer (Cont.)

- **Consistency requirement** in above example:
 - The **sum of A and B** is unchanged by the execution of the transaction
- In general, consistency requirements include the following:
 - Explicitly specified **integrity constraints** such as primary keys and foreign keys
 - During **transaction execution** the database may be **temporarily inconsistent**.
 - When the transaction completes successfully the database must be consistent



Example of Fund Transfer (Cont.)

- **Isolation requirement** — if between steps 3 and 6, another transaction T2 is allowed to access the partially updated database, it will observe an inconsistent database (the sum $A + B$ will be less than that it should be).

T1

1. **read**(A)
2. $A := A - 50$
3. **write**(A)
4. **read**(B)
5. $B := B + 50$
6. **write**(B)

T2

read(A), read(B), write(A+B)

- Isolation can be ensured by running the transactions **serially**.
 - That is, one after the other.
- However, executing multiple transactions concurrently has significant benefits.




Transaction State

- **Active** – it is the initial state; the transaction stays in this state while it is executing
- **Partially committed** – This is the state in which a transaction reaches after the final statement has been executed.
- **Failed** -- A transaction enters the failed state after the system determines that the transaction can no longer proceed with the normal execution.
- **Aborted** – Such a transaction can be rolled back and the database is restored to its state prior to the start of the transaction. At this point the database system has two options:
 - **Restart the transaction**
 - Can be done only if there is no internal logical error
 - **Kill the transaction**
- **Committed** – after successful completion.



Transaction State (Cont.)





Concurrent Executions

- If two transactions are running concurrently, the operating system may execute one transaction for a little while, then perform a **context switch**, execute the second transaction for some time, and then switch back to the first transaction for some time, and so on.
- With multiple transactions, the CPU time is shared among all the transactions.
- In concurrent execution, several execution sequences are possible, since the various instructions from both transactions may now be interleaved.
- In general, it is not possible to predict exactly how many instructions of a transaction will be executed before the CPU switches to another transaction.
- Thus, the number of possible **schedules for a set of n transactions** is larger than **$n!$** .

|



Advantages of Concurrent Executions

- Multiple transactions are allowed to run concurrently in the system.
- The advantages of concurrent execution of transactions are:
 1. **Increased processor and disk utilization**, leading to better transaction *throughput*. For E.g., one transaction can be using the CPU while another is reading from or writing to the disk
 2. **Reduced average response time** for transactions: short transactions need not wait behind long ones.
- **Concurrency control schemes** – Concurrency Control is the management procedure that is required for controlling concurrent execution of the operations that take place on a database.
- The DBS must control the interaction among the concurrent transactions in order to prevent them from destroying the consistency of the database.
- This is done through a variety of mechanisms known as concurrency control schemes.



Schedules of a Transaction

- **Schedule** – a sequences of instructions that specify the chronological order in which instructions of concurrent transactions are executed
 - A schedule for a **set of transactions** must consist of all instructions of those transactions
 - Must preserve the order in which the instructions appear in each individual transaction.
- A transaction that successfully completes its execution will have a **commit instruction** as the last statement
 - By default, a transaction is assumed to execute commit instruction as its last step.
- A transaction that fails to successfully complete its execution will have an **abort instruction** as the last statement



Serial Schedule

- As the name says, all the transactions are executed serially one after the other.
- In serial Schedule, a transaction does not start execution until the currently running transaction finishes execution.

Transaction T1	Transaction T2
R(A)	
W(A)	
R(B)	
W(B)	
commit	
	R(A)
	W(B)
	commit



Schedule 1

- Let T_1 transfer \$50 from A to B , and T_2 transfer 10% of the balance from A to B .
- A **serial** schedule in which T_1 is followed by T_2 :

T_1	T_2
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit



Schedule 2

- A serial schedule where T_2 is followed by T_1

T_1	T_2
	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit	



Concurrent Schedule

- In a concurrent/non-serial schedule, multiple transactions execute concurrently, unlike the serial schedule, where one transaction must wait for another to complete all its operations.
- In this schedule, the other transaction proceeds without the completion of the previous transaction.
- All the transaction operations are interleaved or mixed with each other.

Transaction T1	Transaction T2
R1(A)	
W1(A)	
	R2(A)
	W2(A)
R1(B)	
W1(B)	
	R2(B)
	W2(B)



Schedule 3

- Let T_1 and T_2 be the transactions defined previously. The following schedule is not a serial schedule, but it is *equivalent* to Schedule 1

T_1	T_2
read (A) $A := A - 50$ write (A)	read (A) $temp := A * 0.1$ $A := A - temp$ write (A)
read (B) $B := B + 50$ write (B) commit	read (B) $B := B + temp$ write (B) commit

- In Schedules 1, 2 and 3, the sum $A + B$ is preserved.



- | T_1 | T_2 |
|---|---|
| read (A)
$A := A - 50$ | read (A)
$temp := A * 0.1$
$A := A - temp$
write (A)
read (B) |
| write (A)
read (B)
$B := B + 50$
write (B)
commit | $B := B + temp$
write (B)
commit |



Serializability

- The database system must control concurrent execution of transactions, to ensure that the **database state remains consistent**.
- **Basic Assumption** – Each transaction preserves database consistency.
- Thus, serial execution of a set of transactions preserves database consistency.
- A (possibly concurrent) schedule is **serializable** if it is equivalent to a serial schedule.
- Thus, if a non-serial schedule can be transformed into its corresponding serial schedule, it is said to be **serializable**.
- There are two different forms of serializability in DBMS:
 1. **Conflict Serializability**
 2. **View Serializability**



Conflicting Instructions

- For simplicity, we consider only **read** and **write** instructions in a transaction
- **Conflicting instructions** in a DBMS are operations that are performed on the same data or resource concurrently, and the **order in which they are executed affects the final outcome**
- Conflicting instructions in a transaction occur when two or more transactions try to access the same data item at the same time
- Instructions I_i and I_j of transactions T_i and T_j respectively, **conflict** if and only if there exists some item Q accessed by both I_i and I_j , and at least one of these instructions is a write operation on the same data item.
 1. $I_i = \text{read}(Q)$, $I_j = \text{read}(Q)$. I_i and I_j don't conflict.
 2. $I_i = \text{read}(Q)$, $I_j = \text{write}(Q)$. They conflict.
 3. $I_i = \text{write}(Q)$, $I_j = \text{read}(Q)$. They conflict
 4. $I_i = \text{write}(Q)$, $I_j = \text{write}(Q)$. They conflict
- Thus, I_i and I_j **conflict** if they are operations by different transactions on the same data item, and at least one of these instructions is a write operation



Conflict Serializability

- If a schedule S can be transformed into a schedule S' by a series of swaps of non-conflicting instructions, we say that S and S' are **conflict equivalent** where, S : concurrent schedule, S' : serial schedule.
- In our previous examples, Schedule1 is not conflict equivalent to Schedule2.
- Schedule1 is conflict equivalent to Schedule3, because the non-conflicting instructions of $T1$ - $read(B)$ and $write(B)$ can be swapped with the $read(A)$ and $write(A)$ instructions of $T2$ in Schedule3 to generate its equivalent serial schedule *Schedule1*.
- The concept of conflict equivalence leads to the concept of conflict serializability.
- We say that a non-serial schedule S is **conflict serializable** if it is conflict equivalent to a serial schedule.



Example of Conflict Serializability

T_1	T_2
read(A) $A := A - 50$ write(A) read(B) $B := B + 50$ write(B)	read(A) $temp := A * 0.1$ $A := A - temp$ write(A) read(B) $B := B + temp$ write(B)

T_1	T_2
read(A) $A := A - 50$ write(A)	read(A) $temp := A * 0.1$ $A := A - temp$ write(A)
read(B) $B := B + 50$ write(B)	read(B) $B := B + temp$ write(B)

Schedule 1—a serial schedule in which T_1 is followed by T_2 . Schedule 3—a concurrent schedule equivalent to schedule 1.

Schedule 3 is conflict serializable, since it is conflict equivalent to the serial schedule Schedule 1.



Example of Conflict Serializability (Cont.)

- Schedule 3 can be transformed into Schedule 6, a serial schedule where T_2 follows T_1 , by series of swaps of non-conflicting instructions. Therefore, Schedule 3 is conflict serializable.

T_1	T_2
read (A) write (A)	
	read (A) write (A)
read (B) write (B)	
	read (B) write (B)

Schedule 3

T_1	T_2
read (A) write (A) read (B) write (B)	
	read (A) write (A) read (B) write (B)

Schedule 6



Example of Conflict Serializability (Cont.)

- Example of a schedule that is not conflict serializable:

T_3	T_4
read (Q)	write (Q)
write (Q)	

- The schedule consists of only the significant operations (that is, the read and write) of transactions T_3 and T_4 .
- This schedule is not conflict serializable, since it is not equivalent to either the serial schedule $\langle T_3, T_4 \rangle$ or the serial schedule $\langle T_4, T_3 \rangle$.



View Serializability

- In this form of serializability, we consider a form of equivalence that is less stringent than conflict equivalence, but like conflict equivalence, it is also based on only the read and write operations of transactions.
- Let S and S' be two schedules with the same set of transactions. S and S' are **view equivalent** if the following three conditions are met, for each data item Q :
 1. For each data item Q , if transaction T_i reads the initial value of Q in schedule S , then transaction T_i must, in schedule S' , also read the initial value of Q .
 2. For each data item Q , if transaction T_i executes $\text{read}(Q)$ in schedule S , and if that value was produced by a $\text{write}(Q)$ operation executed by transaction T_j , then the $\text{read}(Q)$ operation of transaction T_i must also read the value of Q that was produced by the same $\text{write}(Q)$ operation of transaction T_j in schedule S' .
 3. For each data item Q , the transaction (if any) that performs the final $\text{write}(Q)$ operation in schedule S must perform the final $\text{write}(Q)$ operation in schedule S' .



View Serializability

Conditions 1 and 2 ensure that each transaction reads the same values in both schedules and, therefore, performs the same computation.

- Condition 3, ensures that both schedules result in the same final system state.
- Schedule 1 is view equivalent to schedule 3, because the values of account *A* and *B* read by transaction *T2* were produced by *T1* in both schedules.
- **View Equivalence:** The concept of view equivalence leads to the concept of view serializability. A schedule *S* is **view serializable if it is view equivalent to a serial schedule**
- Schedule 9 is view serializable. It is view equivalent to the serial schedule $\langle T_3, T_4, T_6 \rangle$, since the one *read(Q)* instruction reads the initial value of *Q* in both schedules, and *T6* performs the final write of *Q* in both the schedules.

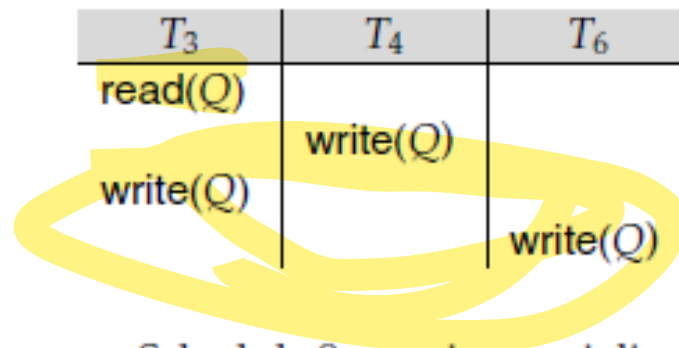
T_3	T_4	T_6
read(<i>Q</i>)	write(<i>Q</i>)	
write(<i>Q</i>)		
		write(<i>Q</i>)

Schedule 9—a view-serializable schedule.



View Serializability (Cont.)

- Every conflict-serializable schedule is also view serializable, but there are view serializable schedules that are not conflict serializable.
- Schedule 9 is not conflict serializable, since every pair of consecutive instructions conflicts, and, thus, no swapping of instructions is possible.
- In Schedule 9, transactions T_4 and T_6 perform $write(Q)$ operations without having performed a $read(Q)$ operation. Writes of this sort are called **blind writes**.
- Blind writes appear in any view-serializable schedule that is not conflict serializable.



Schedule 9—a view-serializable schedule.



Testing for Serializability


- Given a particular schedule S , how to determine, whether the schedule is serializable or not.
- We now present a simple and efficient method for determining conflict serializability of a schedule.
- Consider a schedule S . We construct a directed graph, called a **precedence graph**, from S .
- This graph consists of a pair $G = (V, E)$, where V is a set of **vertices** and E is a set of **edges**.
- The set of vertices consists of all the transactions participating in the schedule.
- The set of edges consists of all edges $T_i \rightarrow T_j$ for which one of three conditions holds:
 1. T_i executes write(Q) before T_j executes read(Q).
 2. T_i executes read(Q) before T_j executes write(Q).
 3. T_i executes write(Q) before T_j executes write(Q).



Testing for Serializability

- If a precedence graph of a Schedule S contains a single edge $T_i \rightarrow T_j$, then all the instructions of T_i are executed before the first instruction of T_j is executed.

Precedence graph for Schedule S



- If a **precedence graph for schedule S contains a cycle**, then **S is non-serializable**.
- If the **precedence graph has no cycle**, then **S is known as serializable**.